

# Introduction to DrCCTProf

Xu Liu

# DrCCTProf Overview

- DrCCTProf functionality
  - ♦ Binary analysis
    - \* Analyze compiler code generation
    - \* Compiler behavior, inefficiencies
  - ♦ Dynamic analysis
    - \* Analyze a program during the program execution with high accuracy
- An open-source project
  - ♦ Hands-on experiences with a real project
    - \* Code with high quality will be committed to the project
  - ♦ Better understand the compiler in real world

# Compilers Do NOT Eliminate All Inefficiencies

- Binary analysis is necessary to check compiler code generation
- Compilers have limitations with their static analysis
  - ♦ Aliasing and pointers
  - ♦ Limited optimization scopes: compilation units
  - ♦ Input-sensitive inefficiencies
  - ♦ Flow-sensitive inefficiencies

# DrCCTProf: Fine-grained Profiling

- Track each instruction
  - ♦ Operator
  - ♦ Operands
- Track each register
  - ♦ General registers
  - ♦ SIMD registers
- Track each memory location
  - ♦ Effective addresses
- Track each value in storage location
  - ♦ Value in registers
  - ♦ Value in memory

# DynamoRIO for Binary Instrumentation

- Robust
  - ♦ Google maintains it
  - ♦ Production-level quality
- Dynamic binary instrumentation
  - ♦ Do not need source code, recompilation, and re-linking.
- Programmable instrumentation
  - ♦ Call-based APIs support to build different instrumentation tools (DynamoRIO clients)
  - ♦ Collecting all aforementioned fine-grained information

# DynamoRIO for Binary Instrumentation

- Robust
  - ♦ Google maintains it
  - ♦ Production-level quality
- Dynamic binary instrumentation
  - ♦ Do not need source code, recompilation, and re-linking.
- Programmable instrumentation
  - ♦ Call-based APIs support to build different instrumentation tools (DynamoRIO clients)
  - ♦ Collecting all aforementioned fine-grained information

**\$ drrun -t client -- application**

# Client Example: Dynamic Basic Block Count

---

```
#include "dr_api.h"
DR_EXPORT void
dr_client_main(client_id_t id, int argc, const char *argv[]) {
    dr_register_bb_event(event_bb);
}
```

```
static void bbcount() { global_count++; }
```

```
static dr_emit_flags_t
event_bb(void *drcontext, void *tag, instrlist_t *bb,
        bool for_trace, bool translating) {
    dr_insert_clean_call(drcontext, bb, instrlist_first_app(bb),
                        (void *)bbcount, false /* save fpstate */, 0);
    return DR_EMIT_DEFAULT;
}
```

# Client Example: Dynamic Basic Block Count

---

```
#include "dr_api.h"
DR_EXPORT void
dr_client_main(client_id_t id, int argc, const char *argv[]) {
    dr_register_bb_event(event_bb);
}
```

```
static void bbcount() { global_count++; }
```



*execution time*

```
static dr_emit_flags_t
event_bb(void *drcontext, void *tag, instrlist_t *bb,
        bool for_trace, bool translating) {
    dr_insert_clean_call(drcontext, bb, instrlist_first_app(bb),
        (void *)bbcount, false /* save fpstate */, 0);
    return DR_EMIT_DEFAULT;
}
```



# Client Example: Dynamic Basic Block Count

```
#include "dr_api.h"
DR_EXPORT void
dr_client_main(client_id_t id, int argc, const char *argv[]) {
    dr_register_bb_event(event_bb);
}
```

```
static void bbcount() { global_count++; }
```



*execution time*

```
static dr_emit_flags_t
event_bb(void *drcontext, void *tag, instrlist_t *bb,
        bool for_trace, bool translating) {
    dr_insert_clean_call(drcontext, bb, instrlist_first_app(bb),
                        (void *)bbcount, false /* save fpstate */, 0);
    return DR_EMIT_DEFAULT;
}
```



*transformation time*

# Client Example: Dynamic Basic Block Count

```
#include "dr_api.h"
DR_EXPORT void
dr_client_main(client_id_t id, int argc, const char *argv[]) {
    dr_register_bb_event(event_bb);
}
```

```
static void bbcount() { global_count++; }
```

```
static dr_emit_flags_t
event_bb(void *drcontext, void *tag, instrlist_t *bb,
        bool for_trace, bool translating) {
    dr_insert_clean_call(drcontext, bb, instrlist_first_
        (void *)bbcount, false /* save
    return DR_EMIT_DEFAULT;
}
```

*execution time*

```
call bbcount()
sub $0xff, %edx
cmp %esi, %edx
jle <L1>
call bbcount()
add %eax, %ebx
mov $0x1, %edi
jmp <L2>
```

*information time*

# DynamoRIO: Instrumentation Granularity

- Instructions
  - ♦ Arithmetic instructions, memory accesses, function calls and returns
- Basic blocks
  - ♦ A sequence of codes ended at a control-flow changing instruction
  - ♦ Single entry, single-exit
- Traces
  - ♦ A sequence of codes ended at an **unconditional** control-flow changing instruction
  - ♦ Single entry, multiple exits

# DynamoRIO: Instrumentation Granularity

- Instructions
  - ♦ Arithmetic instructions, memory accesses, function calls and returns

- Basic blocks
  - ♦ A sequence of codes ended at a control instruction
  - ♦ Single entry, single-exit

```
sub    $0xff, %edx
cmp    %esi, %edx
jle    <L1>
```

- Traces
  - ♦ A sequence of codes ended at an **unconditional** changing instruction
  - ♦ Single entry, multiple exits

```
mov    $0x1, %edi
add    $0x10, %eax
jmp    <L2>
```

1 trace  
2 basic blocks  
6 instructions

# Instrumentation: instruction counting

---

```
DR_EXPORT void
dr_client_main(client_id_t id, int argc, const char *argv[]) {
    ...
    drmgr_register_bb_instrumentation_event
        (event_bb_analysis, event_app_instruction, NULL);
}
```

```
static uint global_count;
static void inscount(uint num_instrs) { global_count += num_instrs; }
```

# Instrumentation: instruction counting

---

```
static dr_emit_flags_t
event_bb_analysis(void *drcontext, void *tag, instrlist_t *bb,
                  bool for_trace, bool translating, void **user_data)
{
    instr_t *instr;
    uint num_instrs;

    for (instr = instrlist_first_app(bb), num_instrs = 0;
         instr != NULL;
         instr = instr_get_next_app(instr)) {
        num_instrs++;
    }

    *user_data = (void *)(ptr_uint_t) num_instrs;

    return DR_EMIT_DEFAULT;
}
```

## Part 3: Instrumentation: instruction counting

---

```
static dr_emit_flags_t
event_app_instruction(void *drcontext, void *tag, instrlist_t *bb, instr_t *instr,
                    bool for_trace, bool translating, void *user_data)
{
    uint num_instrs;

    if (!drmgr_is_first_instr(drcontext, instr))
        return DR_EMIT_DEFAULT;

    num_instrs = (uint)(ptr_uint_t) user_data;
    dr_insert_clean_call(drcontext, bb, instrlist_first_app(bb),
                        (void *)inscount, false /* save fpstate */, 1,
                        OPND_CREATE_INT32(num_instrs));

    return DR_EMIT_DEFAULT;
}
```

# Cross-Platform IR Support

---

- Instruction creation: `XINST_CREATE_*` macros
  - `XINST_CREATE_load()`, `XINST_CREATE_jump()`, etc.
- Generic instruction and operand queries
  - E.g.: `instr_writes_memory()`, `instr_reads_from_reg()`, `instr_compute_address()`, `instr_is_return()`
  - E.g.: `opnd_is_memory_reference()`, `opnd_uses_reg()`
- Generic instrumentation creation helpers
  - E.g.: `instrlist_insert_mov_immed_ptrsz()`, `instrlist_insert_mov_instr_addr()`
- ISA-specific concepts applied to all ISA's
  - E.g., predication
- *drreg* Extension for using registers without naming them



# Cross-Platform IR Support

---

- Instruction creation: `XINST_CREATE_*` macros
  - `XINST_CREATE_load()`, `XINST_CREATE_jump()`, etc.
- Generic instruction and operand queries
  - E.g.: `instr_writes_memory()`, `instr_reads_from_reg()`, `instr_compute_address()`, `instr_is_return()`
  - E.g.: `opnd_is_memory_reference()`, `opnd_uses_reg()`
- Generic instrumentation creation helpers
  - E.g.: `instrlist_insert_mov_immed_ptrsz()`, `instrlist_insert_mov_instr_addr()`
- ISA-specific concepts applied to all ISA's
  - E.g., predication
- *drreg* Extension for using registers without naming them

DynamoRIO works for  
ARM32/64, x86

# Config, Build, and Run

---

- CMake
  - <http://www.cmake.org/>
  - Generates build files for native compiler of choice
    - Makefiles for UNIX, nmake, etc.
    - Visual Studio project files
- CMakeLists.txt

```
add_library(myclient SHARED myclient.c)
find_package(DynamoRIO)
if (NOT DynamoRIO_FOUND)
    message(FATAL_ERROR "DynamoRIO package required to
        build")
endif(NOT DynamoRIO_FOUND)
configure_DynamoRIO_client(myclient)
use_DynamoRIO_extension(myclient drmgr)
```

# Config, Build, and Run

---

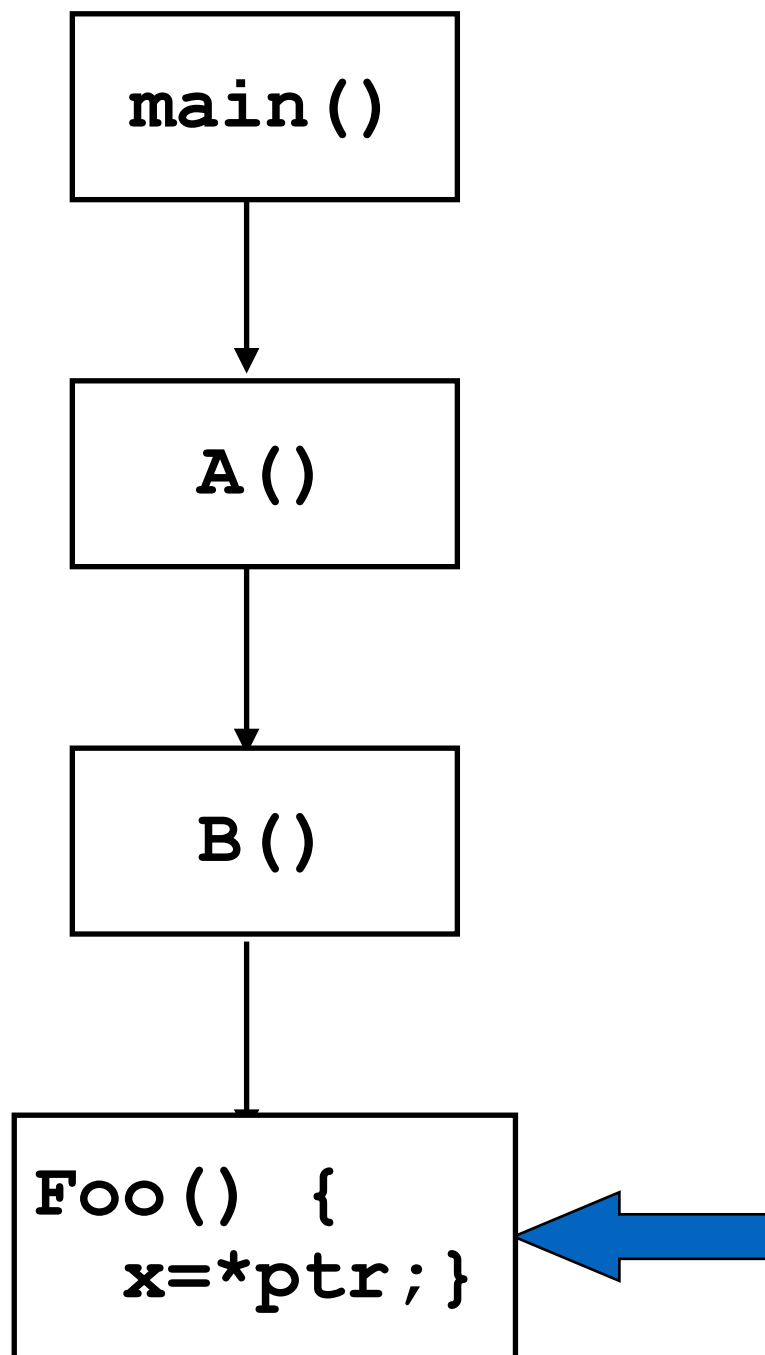
- Config
  - `cmake /path/to/your/client/`
  - `ccmake`, `cmake-gui`
- Build
  - `make`
  - `cmake --build .`
- Run
  - Method 1 (one step)
    - `drrun -c <client> <client options> -- <app cmdline>`
  - Method 2 (two steps, for better child process blacklisting, etc.)
    - `drconfig -reg <appname> -c <client> <client options>`
    - `drinject <app cmdline>`

# Runtime Options

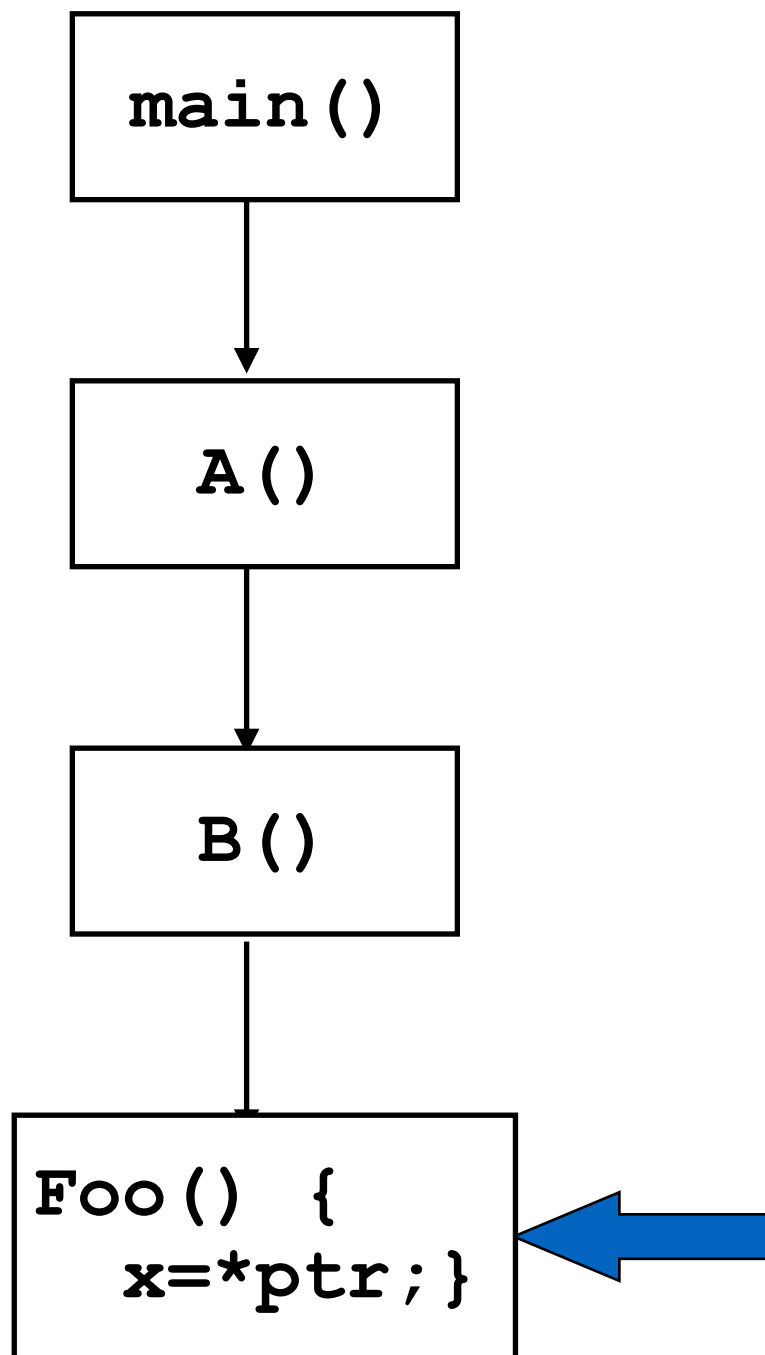
---

- Pass options to drconfig/drrun
- A large number of options; the most relevant are:
  - -t <tool>
  - -c <client lib> <client options>
  - -thread\_private
  - -follow\_children
  - -opt\_cleancall
  - -tracedump\_text and -tracedump\_binary
  - -prof\_pcs
  - -code\_api
  - -help

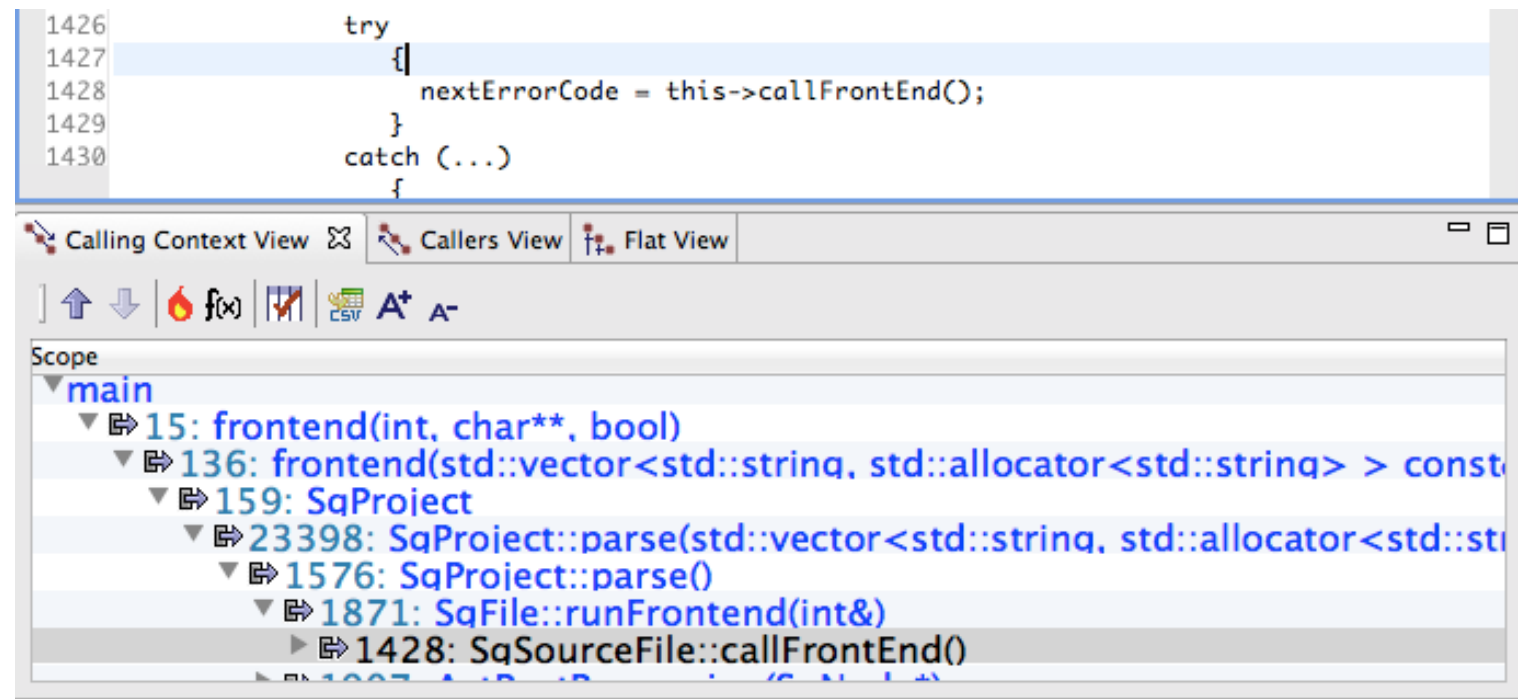
# Collecting Call Path



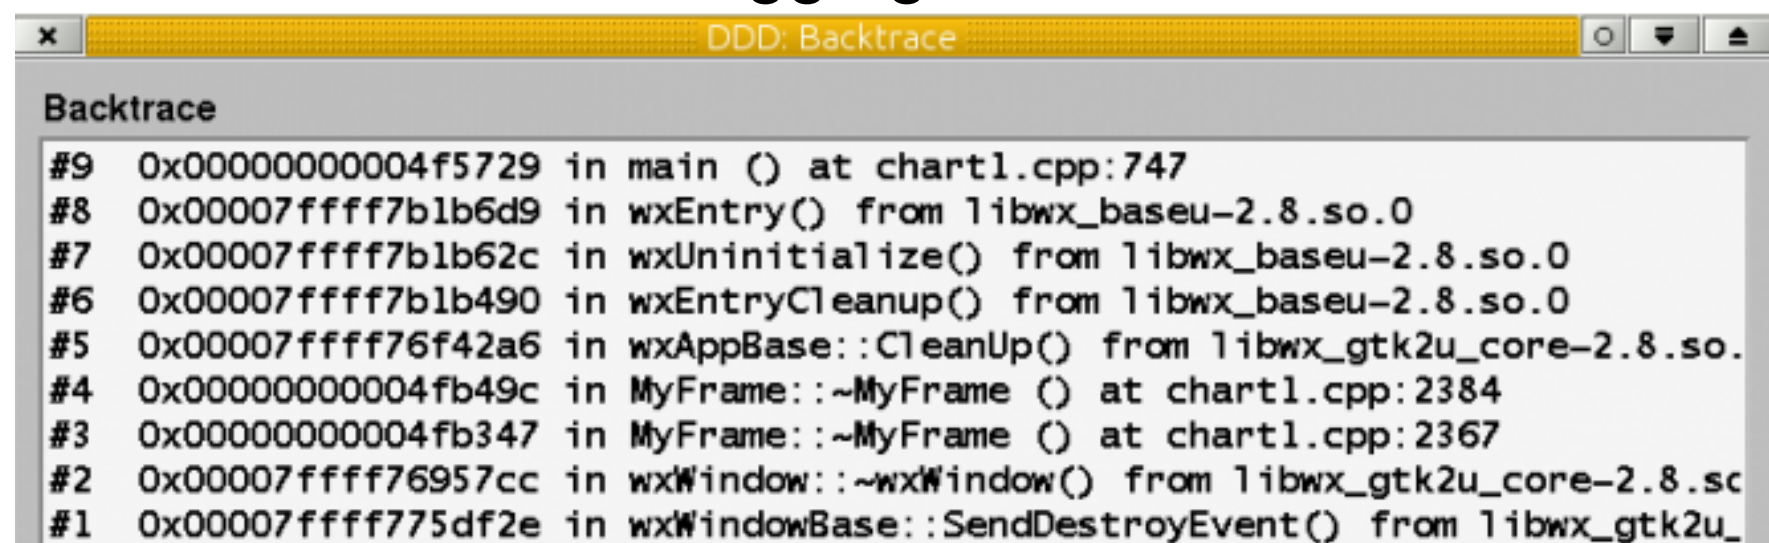
# Collecting Call Path



## Performance Analysis Tools



## Debugging Tools

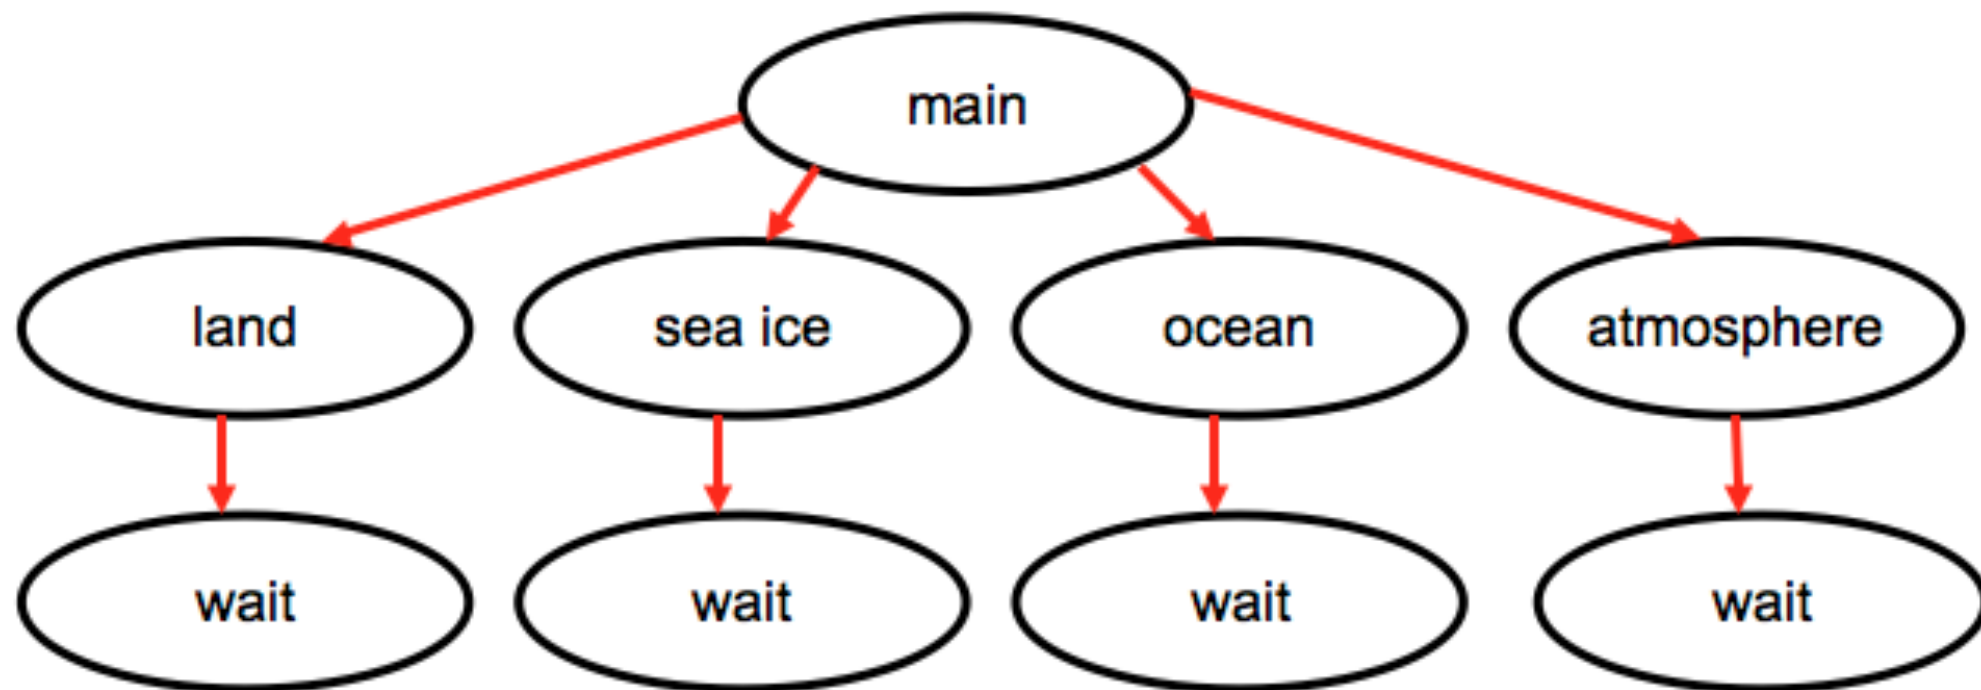


# Call Path Profiling for Fine-grained Analysis

- Associate problematic instructions with their call paths
  - ♦ Expose more semantic information about the instructions
  - ♦ Understand path-sensitive performance issues
- If no call path collected for fine-grained analysis
  - ♦ Do not provide root causes of the problem

# Importance of Call Paths

Example climate code skeleton





# Call Path Profiling for Fine-grained Analysis

- Associate problematic instructions with their call paths
  - ♦ Expose more semantic information about the instructions
  - ♦ Understand path-sensitive performance issues
- If no call path collected for fine-grained analysis
  - ♦ Do not provide root causes of the problem
  - ♦ Do not guide source code optimization

# An Example: SPEC bwaves

A pair of redundant computation

---

```
movsdq 0x8(%rdi,%r10,8), %xmm0:__mul:<no src>
```

\*\*\*\*\*REDUNDANT WITH \*\*\*\*\*

```
movsdq 0x8(%rdi,%r10,8), %xmm0:__mul:<no src>
```

---

# An Example: SPEC bwaves

A pair of redundant computation

---

```
movsdq 0x8(%rdi,%r10,8), %xmm0: __mul:<no src>
__dvd: <no src>
__mpexp: <no src>
__mplog: <no src>
__slowpow: <no src>
__ieee754_pow_sse2: <no src>
pow: <no src>
jacobian_: jacobian_lam.f:47
shell_: shell_lam.f:193
MAIN__: flow_lam.f:63
main: flow_lam.f:67
*****REDUNDANT WITH*****
movsdq 0x8(%rdi,%r10,8), %xmm0: __mul:<no src>
__dvd: <no src>
__mpexp: <no src>
__mplog: <no src>
__slowpow: <no src>
__ieee754_pow_sse2: <no src>
pow: <no src>
jacobian_: jacobian_lam.f:47
shell_: shell_lam.f:193
MAIN__: flow_lam.f:63
main: flow_lam.f:67
```

---

# An Example: SPEC bwaves

A pair of redundant computation

---

```
movsdq 0x8(%rdi,%r10,8), %xmm0: __mul:<no src>
```

```
__dvd: <no src>
```

```
__mpexp: <no src>
```

```
__mplog: <no src>
```

```
__slowpow: <no src>
```

```
__ieee754_pow_sse2: <no src>
```

```
pow: <no src>
```

```
jacobian_: jacobian_lam.f:47
```

```
shell_: shell_lam.f:193
```

```
MAIN__: flow_lam.f:63
```

```
main: flow_lam.f:67
```

\*\*\*\*\*REDUNDANT WITH\*\*\*\*\*

```
movsdq 0x8(%rdi,%r10,8), %xmm0: __mul:<no src>
```

```
__dvd: <no src>
```

```
__mpexp: <no src>
```

```
__mplog: <no src>
```

```
__slowpow: <no src>
```

```
__ieee754_pow_sse2: <no src>
```

```
pow: <no src>
```

```
jacobian_: jacobian_lam.f:47
```

```
shell_: shell_lam.f:193
```

```
MAIN__: flow_lam.f:63
```

```
main: flow_lam.f:67
```

---

# An Example: SPEC bwaves

```
41. ros = q(1, ip1, jp1, kp1)
42. us = q(2, ip1, jp1, kp1)/ros
    ...
47. mu = (mu +
((gm-1.0d0)*(q(5,ip1,jp1,kp1)/ros-
0.5d0*(us*us+vs*vs+ws*ws)))**0.75d0)/
2.0d0
```

A pair of redundant computation

---

```
movsdq 0x8(%rdi,%r10,8), %xmm0: __mul:<no src>
__dvd: <no src>
__mpexp: <no src>
__mplog: <no src>
__slowpow: <no src>
__ieee754_pow_sse2: <no src>
pow: <no src>
jacobian_: jacobian_lam.f:47
shell_: shell_lam.f:193
MAIN__: flow_lam.f:63
main: flow_lam.f:67
*****REDUNDANT WITH *****
movsdq 0x8(%rdi,%r10,8), %xmm0: __mul:<no src>
__dvd: <no src>
__mpexp: <no src>
__mplog: <no src>
__slowpow: <no src>
__ieee754_pow_sse2: <no src>
pow: <no src>
jacobian_: jacobian_lam.f:47
shell_: shell_lam.f:193
MAIN__: flow_lam.f:63
main: flow_lam.f:67
```

---

# An Example: SPEC bwaves

41. ros = q(1, ip1, jp1, kp1)  
42. us = q(2, ip1, jp1, kp1)/ros

...

47. mu = (mu +  
((gm-1.0d0)\*(q(5,ip1,jp1,kp1)/ros-  
0.5d0\*(us\*us+vs\*vs+ws\*ws)))\*\*0.75d0)/  
2.0d0

A pair of redundant computation

---

```
movsdq 0x8(%rdi,%r10,8), %xmm0:__mul:<no src>
```

```
__dvd: <no src>
```

```
__mpexp:<no src>
```

```
__mplog:<no src>
```

```
__slowpow:<no src>
```

```
__ieee754_pow_sse2:<no src>
```

```
pow:<no src>
```

```
jacobian_:jacobian_lam.f:47
```

```
shell_:shell_lam.f:193
```

```
MAIN__:flow_lam.f:63
```

```
main:flow_lam.f:67
```

\*\*\*\*\*REDUNDANT WITH \*\*\*\*\*

```
movsdq 0x8(%rdi,%r10,8), %xmm0:__mul:<no src>
```

```
__dvd:<no src>
```

```
__mpexp:<no src>
```

```
__mplog:<no src>
```

```
__slowpow:<no src>
```

```
__ieee754_pow_sse2:<no src>
```

```
pow:<no src>
```

```
jacobian_:jacobian_lam.f:47
```

```
shell_:shell_lam.f:193
```

```
MAIN__:flow_lam.f:63
```

```
main:flow_lam.f:67
```

---

# An Example: SPEC bwaves

41. ros = q(1, ip1, jp1, kp1)  
42. us = q(2, ip1, jp1, kp1)/ros

...

47. mu = (mu +  
((gm-1.0d0)\*(q(5,ip1,jp1,kp1)/ros-  
0.5d0\*(us\*us+vs\*vs+ws\*ws)))\*\*0.75d0)/  
2.0d0

No insights without call path  
profiling

A pair of redundant computation

---

```
movsdq 0x8(%rdi,%r10,8), %xmm0:__mul:<no src>
__dvd:<no src>
__mpexp:<no src>
__mplog:<no src>
__slowpow:<no src>
__ieee754_pow_sse2:<no src>
pow:<no src>
jacobian_:jacobian_lam.f:47
shell_:shell_lam.f:193
MAIN__:flow_lam.f:63
main:flow_lam.f:67
*****REDUNDANT WITH *****
movsdq 0x8(%rdi,%r10,8), %xmm0:__mul:<no src>
__dvd:<no src>
__mpexp:<no src>
__mplog:<no src>
__slowpow:<no src>
__ieee754_pow_sse2:<no src>
pow:<no src>
jacobian_:jacobian_lam.f:47
shell_:shell_lam.f:193
MAIN__:flow_lam.f:63
main:flow_lam.f:67
```

---

# An Example: SPEC bwaves

```
41. ros = q(1, ip1, jp1, kp1)
42. us = q(2, ip1, jp1, kp1)/ros
```

...

```
47. mu = (mu +
((gm-1.0d0)*(q(5,ip1,jp1,kp1)/ros-
0.5d0*(us*us+vs*vs+ws*ws)))**0.75d0)/
2.0d0
```

No insights without call path  
profiling

A pair of redundant computation

---

```
movsdq 0x8(%rdi,%r10,8), %xmm0: __mul:<no src>
__dvd: <no src>
__mpexp:<no src>
__mplog:<no src>
__slowpow:<no src>
__ieee754_pow_sse2:<no src>
pow:<no src>
jacobian_:jacobian_lam.f:47
shell_:shell_lam.f:193
MAIN__:flow_lam.f:63
main:flow_lam.f:67
*****REDUNDANT WITH *****
movsdq 0x8(%rdi,%r10,8), %xmm0: __mul:<no src>
__dvd:<no src>
__mpexp:<no src>
__mplog:<no src>
__slowpow:<no src>
__ieee754_pow_sse2:<no src>
```

DrCCTProf: a framework that collects calling context for  
fine-grained profilers



# DrCCTProf Overview

- Functionality
  - ♦ capture call path for each dynamic instruction
  - ♦ capture the data object read/written by each memory access
    - \* heap data objects: call paths to the allocations
    - \* static data objects: names from symbol table
- Programmability
  - ♦ APIs provide request-based service for clients
- Overhead
  - ♦ moderate overhead in both runtime and space

# DrCCTProf Tutorial Outline

software download

<https://github.com/Xuhpclab/DrCCTProf>

- DrCCTProf APIs
- General usage of DrCCTProf framework for fine-grained profilers
- DrCCTProf framework internals

# DrCCTProf Tutorial Outline

software download

<https://github.com/Xuhpclab/DrCCTProf>

- DrCCTProf APIs
- General usage of DrCCTProf framework for fine-grained profilers
- DrCCTProf framework internals

# A Simple DrCCTProf Client

```
#include <iterator>
#include "dr_api.h"
#include "drcctlb.h"
FILE* gTraceFile;
```

```
void InsertCleansall(int32_t slot) {
    void *drcontext = dr_get_current_drcontext();
    context_handle_t cur_ctxt_hdl = drcctlb_get_context_handle(drcontext, slot);
}

void InstrumentIns(void *drcontext, instr_instrument_msg_t *instrument_msg) {
    dr_insert_clean_call(drcontext, instrument_msg->bb, instrument_msg->instr,
        (void *)InsertCleansall, false, 1, OPND_CREATE_INT32(instrument_msg->slot));
}

void ClientInit(int argc, const char *argv[]) {

}

void ClientExit(void) {
    drcctlb_exit();
}

void dr_client_main(client_id_t id, int argc, const char *argv[]) {
    dr_set_client_name("DynamoRIO Client 'drcctlb_cct_only_clean_call'",
        "http://dynamorio.org/issues");
    ClientInit(argc, argv);

    drcctlb_init(INTERESTING_INS_ALL, gTraceFile, InstrumentIns, false);

    dr_register_exit_event(ClientExit);
}
```

# A Simple DrCCTProf Client

```
#include <iterator>
#include "dr_api.h"
#include "drcctlb.h"
FILE* gTraceFile;
```

```
void InsertCleancall(int32_t slot) {
    void *drcontext = dr_get_current_drcontext();
    context_handle_t cur_ctxt_hdl = drcctlb_get_context_handle(drcontext, slot);
}

void InstrumentIns(void *drcontext, instr_instrument_msg_t *instrument_msg) {
    dr_insert_clean_call(drcontext, instrument_msg->bb, instrument_msg->instr,
        (void *)InsertCleancall, false, 1, OPND_CREATE_INT32(instrument_msg->slot));
}

void ClientInit(int argc, const char *argv[]) {
}

void ClientExit(void) {
    drcctlb_exit();
}

void dr_client_main(client_id_t id, int argc, const char *argv[]) {
    dr_set_client_name("DynamoRIO Client 'drcctlb_cct_only_clean_call'",
        "http://dynamorio.org/issues");
    ClientInit(argc, argv);

    drcctlb_init(INTERESTING_INS_ALL, gTraceFile, InstrumentIns, false);

    dr_register_exit_event(ClientExit);
}
```

# Functional APIs

```
drcctlb_init(bool (*IsInterestingIns)(instr_t *), file_t logFile,  
            void (*userCallBack)(void *, instr_instrument_msg_t *),  
            bool doDataCentric);
```

## ♦ Description:

- \* DrCCTProf clients must call this before using DrCCTProf

## ♦ Arguments:

- \* IsInterestingIns: tells whether a given INS needs to collect context;  
predefined: INTERESTING\_INS\_ALL, INTERESTING\_INS\_NONE,  
INTERESTING\_INS\_MEMORY\_ACCESS.
- \* logFile: file pointer where DrCCTProf put output data
- \* userCallBack: a client callback on each INS that IsInterestingIns is true
- \* doDataCentric: should be set if the client wants DrCCTProf to do data-centric attribution

# A Simple CCTLib Client

```
#include <iterator>
#include "dr_api.h"
#include "drcctlb.h"
FILE* gTraceFile;

void InsertCleancall(int32_t slot) {
    void *drcontext = dr_get_current_drcontext();
    context_handle_t cur_ctxt_hndl = drcctlb_get_context_handle(drcontext, slot);
}

void InstrumentIns(void *drcontext, instr_instrument_msg_t *instrument_msg) {
    dr_insert_clean_call(drcontext, instrument_msg->bb, instrument_msg->instr,
        (void *)InsertCleancall, false, 1, OPND_CREATE_INT32(instrument_msg->slot));
}

void ClientInit(int argc, const char *argv[]) {

}

void ClientExit(void) {
    drcctlb_exit();
}

void dr_client_main(client_id_t id, int argc, const char *argv[]) {
    dr_set_client_name("DynamoRIO Client 'drcctlb_cct_only_clean_call'",
        "http://dynamorio.org/issues");
    ClientInit(argc, argv);
    drcctlb_init(INTERESTING_INS_ALL, gTraceFile, InstrumentIns, false);

    dr_register_exit_event(ClientExit);
}
```

# A Simple CCTLib Client

```
#include <iterator>
#include "dr_api.h"
#include "drcctlb.h"
FILE* gTraceFile;

void InsertCnecall(int32_t slot) {
    void *drcontext = dr_get_current_drcontext();
    context_handle_t cur_ctxt_hndl = drcctlb_get_context_handle(drcontext, slot);
}

void InstrumentIns(void *drcontext, instr_instrument_msg_t *instrument_msg) {
    dr_insert_clean_call(drcontext, instrument_msg->bb, instrument_msg->instr,
        (void *)InsertCnecall, false, 1, OPND_CREATE_INT32(instrument_msg->slot));
}

void ClientInit(int argc, const char *argv[]) {

}

void ClientExit(void) {
    drcctlb_exit();
}

void dr_client_main(client_id_t id, int argc, const char *argv[]) {
    dr_set_client_name("DynamoRIO Client 'drcctlb_cct_only_clean_call'",
        "http://dynamorio.org/issues");
    ClientInit(argc, argv);
    drcctlb_init(INTERESTING_INS_ALL, gTraceFile, InstrumentIns, false);
    monitor every instruction
    dr_register_exit_event(ClientExit);
}
```



# A Simple CCTLib Client

```
#include <iterator>
#include "dr_api.h"
#include "drcctlb.h"
FILE* gTraceFile;

void InsertCancall(int32_t slot) {
    void *drcontext = dr_get_current_drcontext();
    context_handle_t cur_ctxt_hndl = drcctlb_get_context_handle(drcontext, slot);
}

void InstrumentIns(void *drcontext, instr_instrument_msg_t *instrument_msg) {
    dr_insert_clean_call(drcontext, instrument_msg->bb, instrument_msg->instr,
        (void *)InsertCancall, false, 1, OPND_CREATE_INT32(instrument_msg->slot));
}

void ClientInit(int argc, const char *argv[]) {
}

void ClientExit(void) {
    drcctlb_exit();
}

void dr_client_main(client_id_t id, int argc, const char *argv[]) {
    dr_set_client_name("DynamoRIO Client 'drcctlb_cct_only_clean_call'",
        "http://dynamorio.org/issues");
    ClientInit(argc, argv);
    drcctlb_init(INTERESTING_INS_ALL, gTraceFile, InstrumentIns, false);
    monitor every instruction
    dr_register_exit_event(ClientExit);
}
```

# A Simple CCTLib Client

```
#include <iterator>
#include "dr_api.h"
#include "drcctlb.h"
FILE* gTraceFile;

void InsertCancall(int32_t slot) {
    void *drcontext = dr_get_current_drcontext();
    context_handle_t cur_ctxt_hndl = drcctlb_get_context_handle(drcontext, slot);
}

void InstrumentIns(void *drcontext, instr_instrument_msg_t *instrument_msg) {
    dr_insert_clean_call(drcontext, instrument_msg->bb, instrument_msg->instr,
        (void *)InsertCancall, false, 1, OPND_CREATE_INT32(instrument_msg->slot));
}

void ClientInit(int argc, const char *argv[]) {
}

void ClientExit(void) {
    drcctlb_exit();
}

void dr_client_main(client_id_t id, int argc, const char *argv[]) {
    dr_set_client_name("DynamoRIO Client 'drcctlb_cct_only_clean_call'",
        "http://dynamorio.org/issues");
    ClientInit(argc, argv);
    drcctlb_init(INTERESTING_INS_ALL, gTraceFile, InstrumentIns, false);
    dr_register_exit_event(ClientExit);
}
```

The diagram illustrates the flow of control between the client functions and the main function. A green dashed arrow originates from the `gTraceFile` variable in the `drcctlb_init` call within `dr_client_main` and points to its declaration at the top of the file. A blue dashed arrow originates from the `InstrumentIns` function call within `drcctlb_init` and points to its definition. A red dashed arrow originates from the `ClientExit` function call within `dr_client_main` and points to its definition.

monitor every instruction

# A Simple CCTLib Client

```
#include <iterator>
#include "dr_api.h"
#include "drcctlb.h"
FILE* gTraceFile;
```

```
void InsertCancall(int32_t slot) {
    void *drcontext = dr_get_current_drcontext();
    context_handle_t cur_ctxt_hndl = drcctlb_get_context_handle(drcontext, slot);
}
```

```
void InstrumentIns(void *drcontext, instr_instrument_msg_t *instrument_msg) {
    dr_insert_clean_call(drcontext, instrument_msg->bb, instrument_msg->instr,
        (void *)InsertCancall, false, 1, OPND_CREATE_INT32(instrument_msg->slot));
}
```

instrumentation routine

```
void ClientInit(int argc, const char *argv[]) {
}
```

```
void ClientExit(void) {
    drcctlb_exit();
}
```

```
void dr_client_main(client_id_t id, int argc, const char *argv[]) {
    dr_set_client_name("DynamoRIO Client 'drcctlb_cct_only_clean_call'",
        "http://dynamorio.org/issues");
    ClientInit(argc, argv);
    drcctlb_init(INTERESTING_INS_ALL, gTraceFile, InstrumentIns, false);
    monitor every instruction
    dr_register_exit_event(ClientExit);
}
```

# A Simple CCTLib Client

```
#include <iterator>
#include "dr_api.h"
#include "drcctlb.h"
FILE* gTraceFile;
```

```
void InsertCancall(int32_t slot) {
    void *drcontext = dr_get_current_drcontext();
    context_handle_t cur_ctxt_hdl = drcctlb_get_context_handle(drcontext, slot);
}
```

```
void InstrumentIns(void *drcontext, instr_instrument_msg_t *instrument_msg) {
    dr_insert_clean_call(drcontext, instrument_msg->bb, instrument_msg->instr,
        (void *)InsertCancall, false, 1, OPND_CREATE_INT32(instrument_msg->slot));
}
```

instrumentation routine

```
void ClientInit(int argc, const char *argv[]) {
}
```

```
void ClientExit(void) {
    drcctlb_exit();
}
```

```
void dr_client_main(client_id_t id, int argc, const char *argv[]) {
    dr_set_client_name("DynamoRIO Client 'drcctlb_cct_only_clean_call'",
        "http://dynamorio.org/issues");
    ClientInit(argc, argv);
    drcctlb_init(INTERESTING_INS_ALL, gTraceFile, InstrumentIns, false);
    dr_register_exit_event(ClientExit);
}
```

monitor every instruction

no data-centric

# A Simple CCTLib Client

```
#include <iterator>
#include "dr_api.h"
#include "drcctlb.h"
FILE* gTraceFile;
```

```
void InsertCancall(int32_t slot) {
    void *drcontext = dr_get_current_drcontext();
    context_handle_t cur_ctxt_hdl = drcctlb_get_context_handle(drcontext, slot);
}
```

```
void InstrumentIns(void *drcontext, instr_instrument_msg_t *instrument_msg) {
    dr_insert_clean_call(drcontext, instrument_msg->bb, instrument_msg->instr,
        (void *)InsertCancall, false, 1, OPND_CREATE_INT32(instrument_msg->slot));
}
```

instrumentation routine

```
void ClientInit(int argc, const char *argv[]) {
}
```

```
void ClientExit(void) {
    drcctlb_exit();
}
```

```
void dr_client_main(client_id_t id, int argc, const char *argv[]) {
    dr_set_client_name("DynamoRIO Client 'drcctlb_cct_only_clean_call'",
        "http://dynamorio.org/issues");
    ClientInit(argc, argv);
    drcctlb_init(INTERESTING_INS_ALL, gTraceFile, InstrumentIns, false);
    dr_register_exit_event(ClientExit);
}
```

monitor every instruction

no data-centric

# A Simple CCTLib Client

```
#include <iterator>
#include "dr_api.h"
#include "drcctlb.h"
FILE* gTraceFile;
```

```
void InsertCancall(int32_t slot) {
    void *drcontext = dr_get_current_drcontext();
    context_handle_t cur_ctxt_hndl = drcctlb_get_context_handle(drcontext, slot);
}
```

```
void InstrumentIns(void *drcontext, instr_instrument_msg_t *instrument_msg) {
    dr_insert_clean_call(drcontext, instrument_msg->bb, instrument_msg->instr,
        (void *)InsertCancall, false, 1, OPND_CREATE_INT32(instrument_msg->slot));
}
```

instrumentation routine

```
void ClientInit(int argc, const char *argv[]) {
}
```

```
void ClientExit(void) {
    drcctlb_exit();
}
```

```
void dr_client_main(client_id_t id, int argc, const char *argv[]) {
    dr_set_client_name("DynamoRIO Client 'drcctlb_cct_only_clean_call'",
        "http://dynamorio.org/issues");
    ClientInit(argc, argv);
    drcctlb_init(INTERESTING_INS_ALL, gTraceFile, InstrumentIns, false);
    dr_register_exit_event(ClientExit);
}
```

monitor every instruction

no data-centric

# A Simple CCTLib Client

```
#include <iterator>
#include "dr_api.h"
#include "drcctlb.h"
FILE* gTraceFile;
```

```
void InsertCancall(int32_t slot) {
    void *drcontext = dr_get_current_drcontext();
    context_handle_t cur_ctxt_hdl = drcctlb_get_context_handle(drcontext, slot);
}
```

analysis routine

```
void InstrumentIns(void *drcontext, instr_instrument_msg_t *instrument_msg) {
    dr_insert_clean_call(drcontext, instrument_msg->bb, instrument_msg->instr,
        (void *)InsertCancall, false, 1, OPND_CREATE_INT32(instrument_msg->slot));
}
```

instrumentation routine

```
void ClientInit(int argc, const char *argv[]) {
}
```

```
void ClientExit(void) {
    drcctlb_exit();
}
```

```
void dr_client_main(client_id_t id, int argc, const char *argv[]) {
    dr_set_client_name("DynamoRIO Client 'drcctlb_cct_only_clean_call'",
        "http://dynamorio.org/issues");
    ClientInit(argc, argv);
    drcctlb_init(INTERESTING_INS_ALL, gTraceFile, InstrumentIns, false);
    dr_register_exit_event(ClientExit);
}
```

monitor every instruction

no data-centric



# Functional APIs

- **context\_handle\_t drcctlib\_get\_context\_handle**(void \*drcontext, int32\_t opaqueHandle);
  - ♦ Description:
    - \* Get the calling context handle (**context\_handle\_t**)
  - ♦ Arguments:
    - \* drcontext: Dynamorio's thread private context of the asking thread.
    - \* opaqueHandle: handle passed by DrCCTProf to the client tool in its userCallback.
- **data\_handle\_t drcctlib\_get\_data\_hndl**(void \*drcontext, void \*address);
  - ♦ Description:
    - \* Call when need the handle to a data object (**data\_handle\_t**)
  - ♦ Arguments:
    - \* drcontext: Dynamorio's thread private context of the asking thread.
    - \* address: effective address for which the data object is needed.



# Data Structures

- **context\_handle\_t;**
  - ♦ serve as ID of a context
- **data\_handle\_t**
  - ♦ enum objectType:
    - \* STACK\_OBJECT, DYNAMIC\_OBJECT, STATIC\_OBJECT, UNKNOWN
  - ♦ union {path\_handle, sym\_name}:
    - \* path\_handle represents the allocation point of dynamic data
    - \* sym\_name represents the name of static data
  - ♦ uint64\_t beg\_addr
  - ♦ uint64\_t end\_addr
  - ♦ Used in data-centric analysis

# Functional APIs

- `context_t * drcctlib_get_full_cct(context_handle_t ctxtHandle, int maxDepth);`
  - ♦ Description:
    - \* Return the full calling context list whose handle is ctxtHandle
  - ♦ Arguments:
    - \* ctxtHandle: the context handle which request the call path
    - \* maxDepth: the max depth of the returned calling context
- `void drcctlib_print_full_cct(context_handle_t ctxtHandle, int maxDepth);`
  - ♦ Description:
    - \* Prints the full calling context whose handle is ctxtHandle
    - \* Often used when log the analysis result to file
  - ♦ Arguments:
    - \* ctxtHandle: the context handle which request the call path
    - \* maxDepth: the max depth of the printed calling context

# Functional APIs

- **char \*** **drcctlib\_get\_str\_from\_strpool**(int32\_t index);
  - ♦ Description:
    - \* Get the char string for a symbol from string pool index
    - \* Often used to get the name of Static data objects(pass sym\_name)
  - ♦ Arguments:
    - \* index: a string pool index
- **bool drcctlib\_have\_same\_caller\_prefix**(context\_handle\_t ctxt1, context\_handle\_t ctxt2);
  - ♦ Description:
    - \* Tell if the call path from root to leaves except the leaves themselves are the same for the given two context handles
- **bool drcctlib\_have\_same\_source\_line**(context\_handle\_t ctxt1, context\_handle\_t ctxt2);
  - ♦ Description:
    - \* Tell if the two given context handles point to the same source line, no matter whether the full call path is same or not

# DrCCTProf Tutorial Outline

software download

<https://github.com/Xuhpclab/DrCCTProf>

- DrCCTProf APIs
- General usage of DrCCTProf framework for fine-grained profilers
- DrCCTProf framework internals

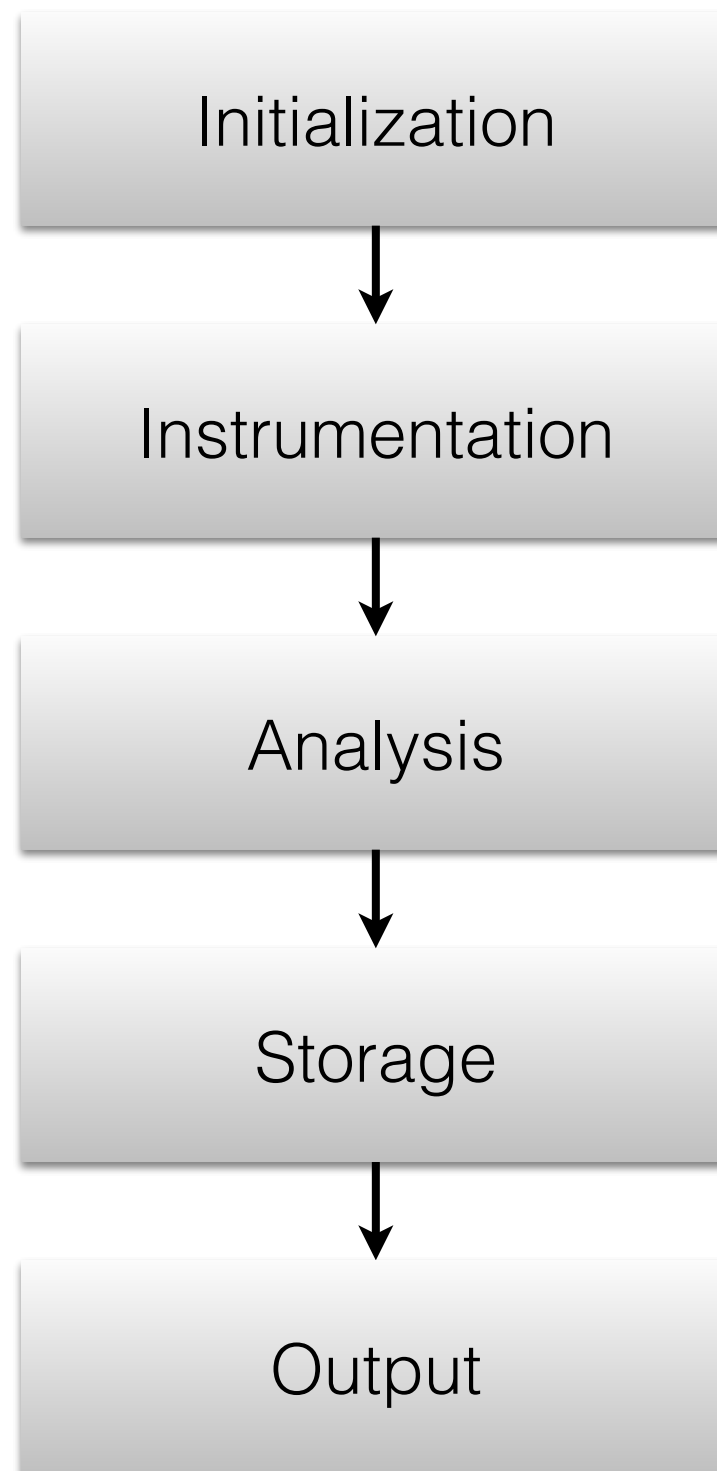
# DrCCTProf Tutorial Outline

software download

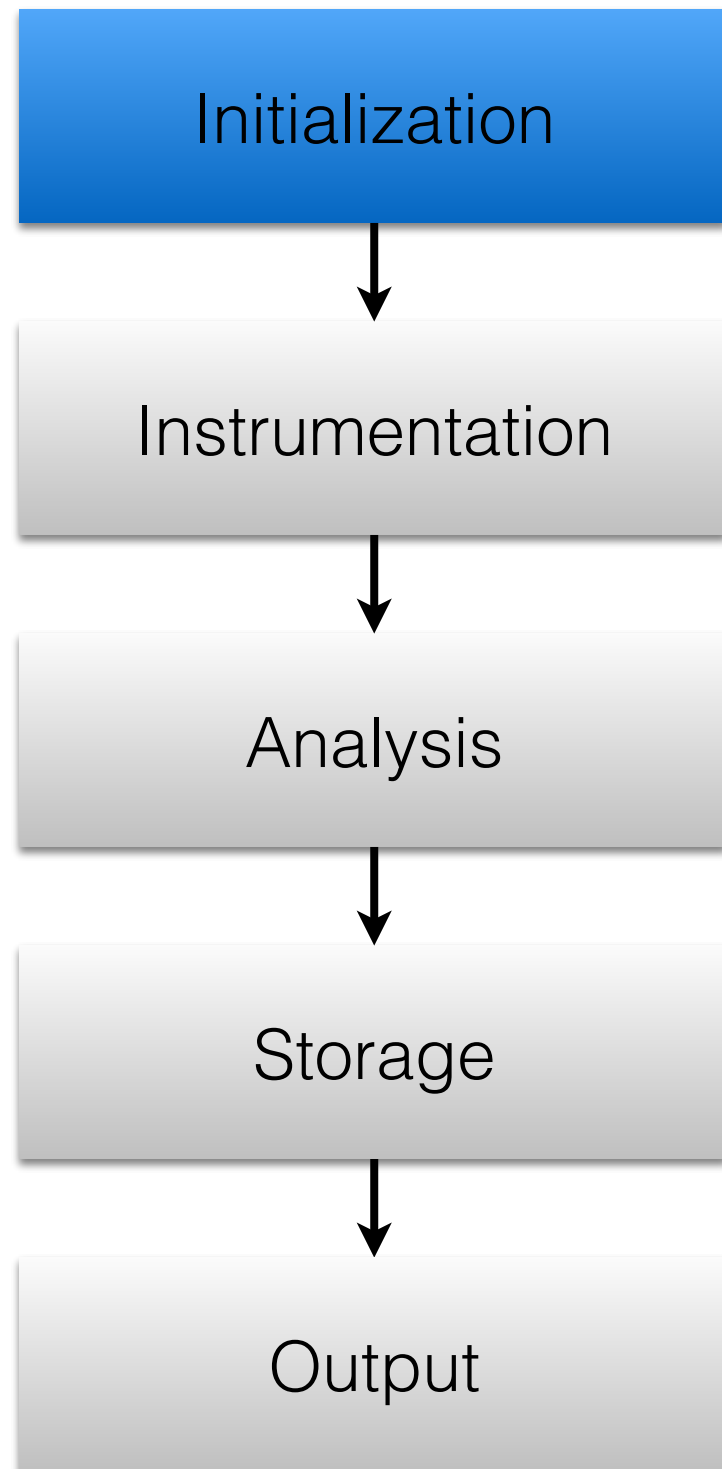
<https://github.com/Xuhpclab/DrCCTProf>

- DrCCTProf APIs
- General usage of DrCCTProf framework for fine-grained profilers
- DrCCTProf framework internals

# Work Flow to Prepare DrCCTProf Clients

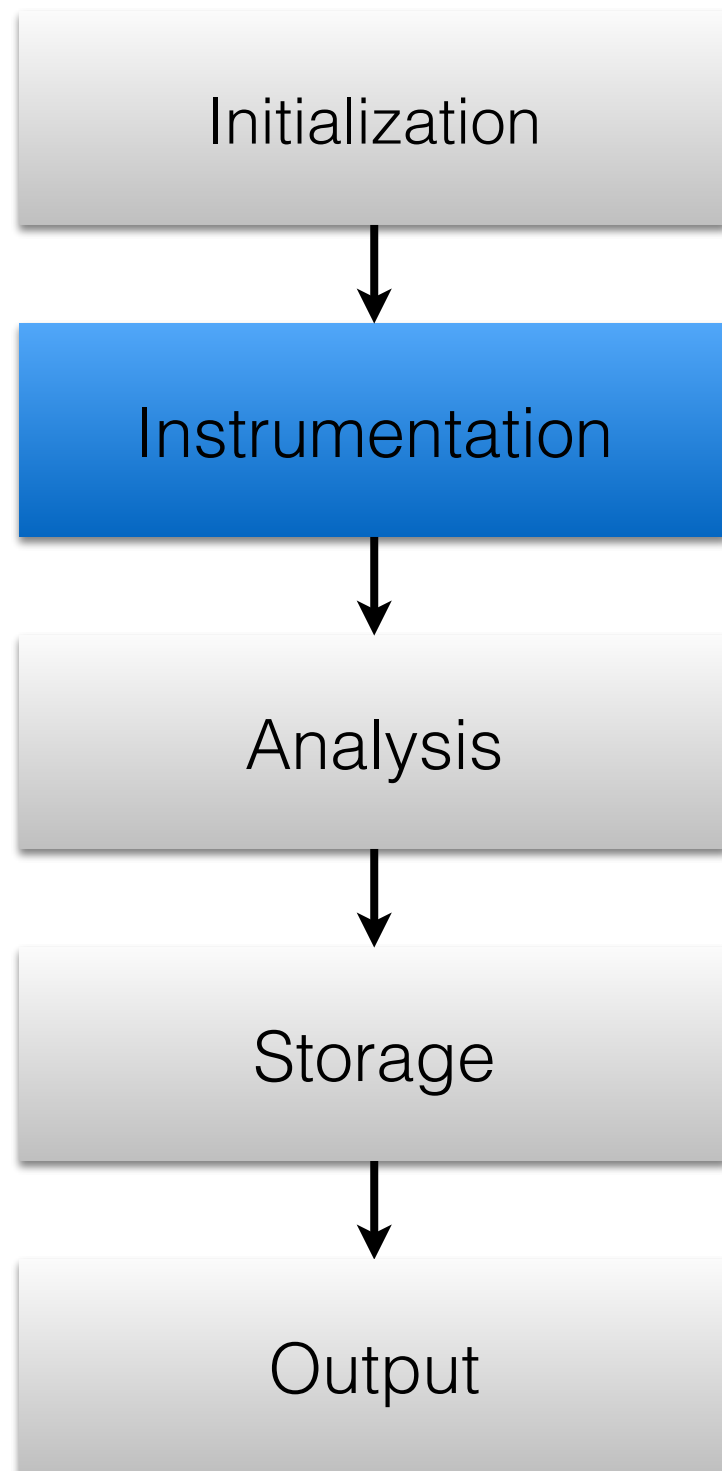


# Work Flow to Prepare DrCCTProf Clients



- Initialize Dynamorio components required by the client
  - ♦ `drmgr_init();`
  - ♦ `drsym_init();`
  - ♦ `drwrap_init();`
  - ♦ ...
- Initialize DrCCTProf and client
  - ♦ Create output file for client
  - ♦ **`drcctlib_init()`**
    - \* Initial instruction to focus
    - \* Enable Data-centric or not
- Initialize data structures
  - ♦ Obtain key for TLS
    - \* `drmgr_register_tls_field()`
  - ♦ Define & initialize data items
    - \* Define a structure with elements you want to store per thread

# Work Flow to Prepare DrCCTProf Clients

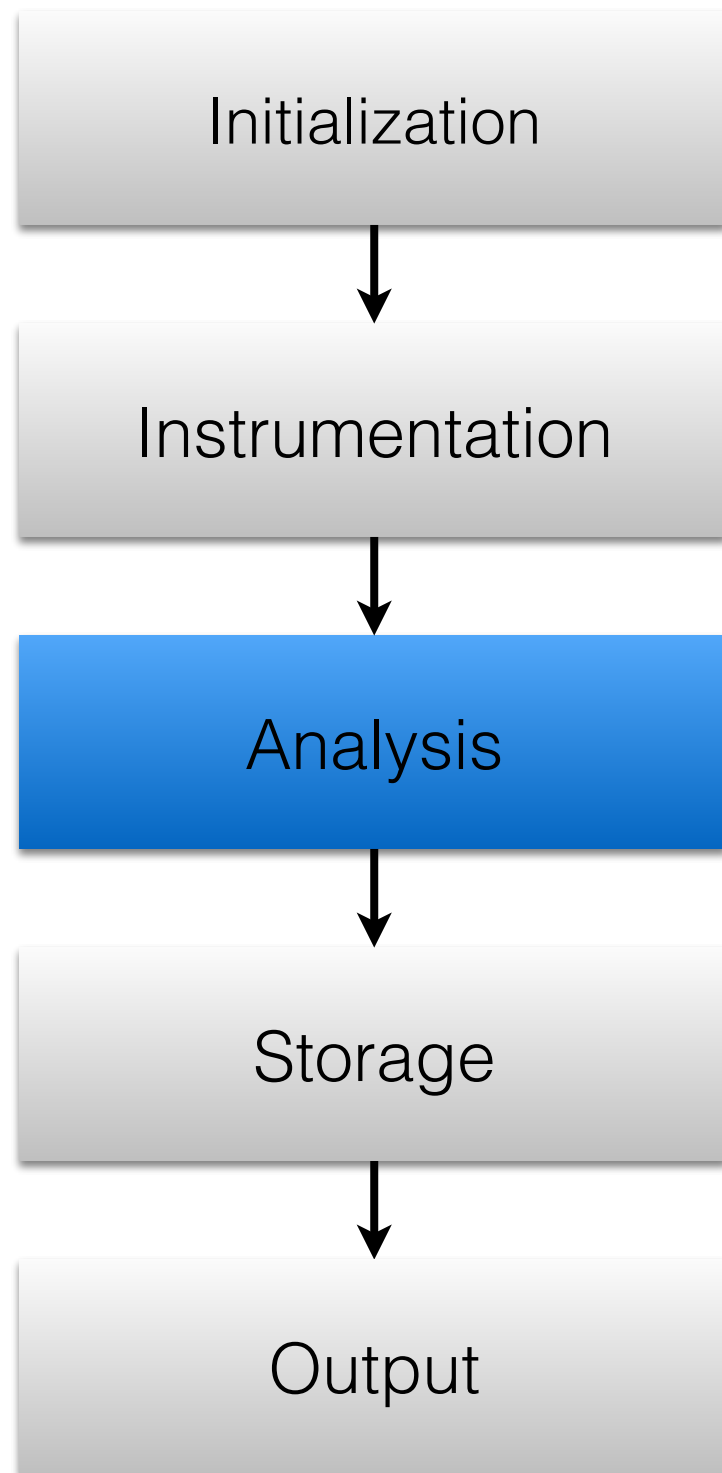


- Instrumentation APIs
  - ♦ `drmgr_register_bb_instrumentation_event()`
  - ♦ `drmgr_register_module_load_event()`
  - ♦ `drwrap_wrap()`
  - ♦ ....
- INS insert calls
  - ♦ `drcontext`: Dynamorio's thread private context of the asking thread.
  - ♦ `internal_instrument_msg`: contain Dynamorio's basic block struct, Dynamorio's instruction struct, and opaque handle passed by DrCCTProf

```
InstrumentIns(void *drcontext, instr_instrument_msg_t *internal_instrument_msg)
{
    .....
    .....
    dr_insert_clean_call(drcontext,
        internal_instrument_msg->bb, internal_instrument_msg->instr,
        (void *)InsertCleanscall, false, 1,
        OPND CREATE INT32(internal_instrument_msg->slot));
}
```

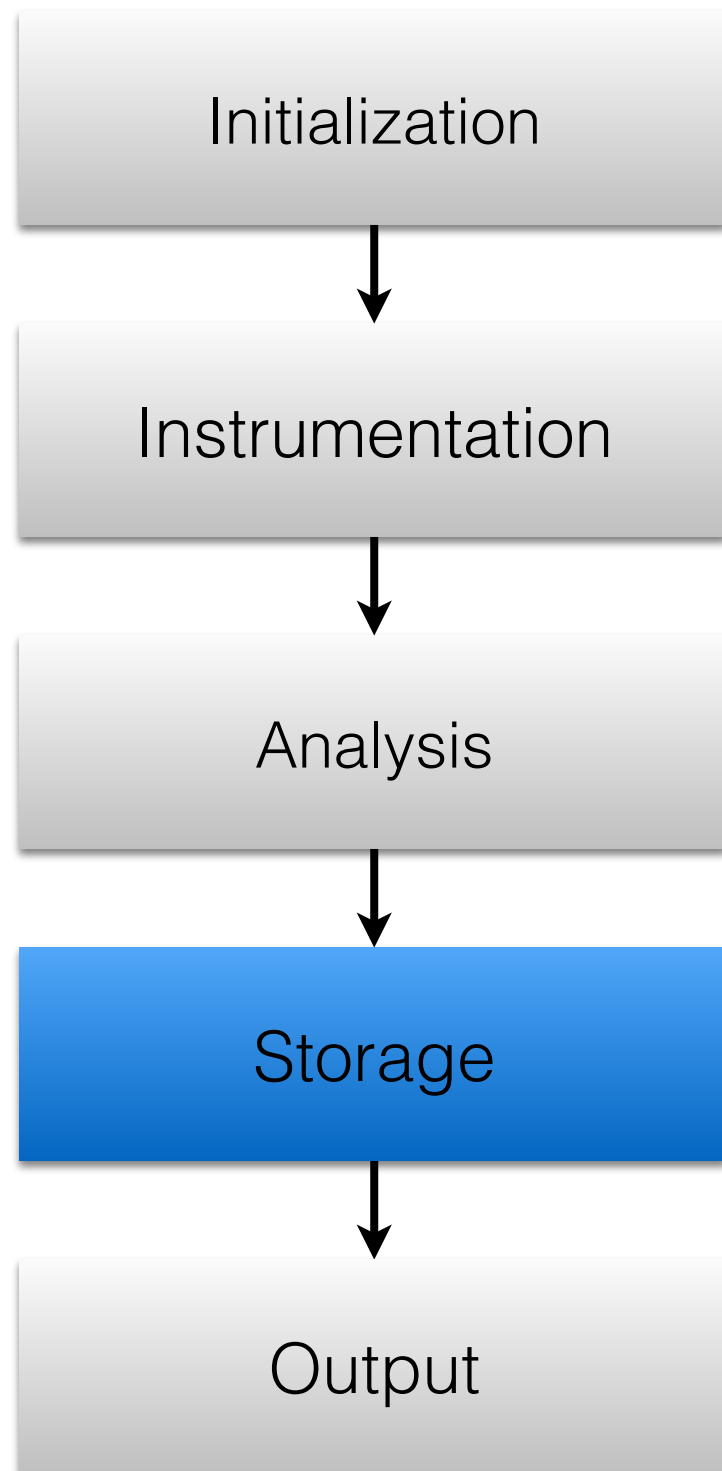


# Work Flow to Prepare DrCCTProf Clients



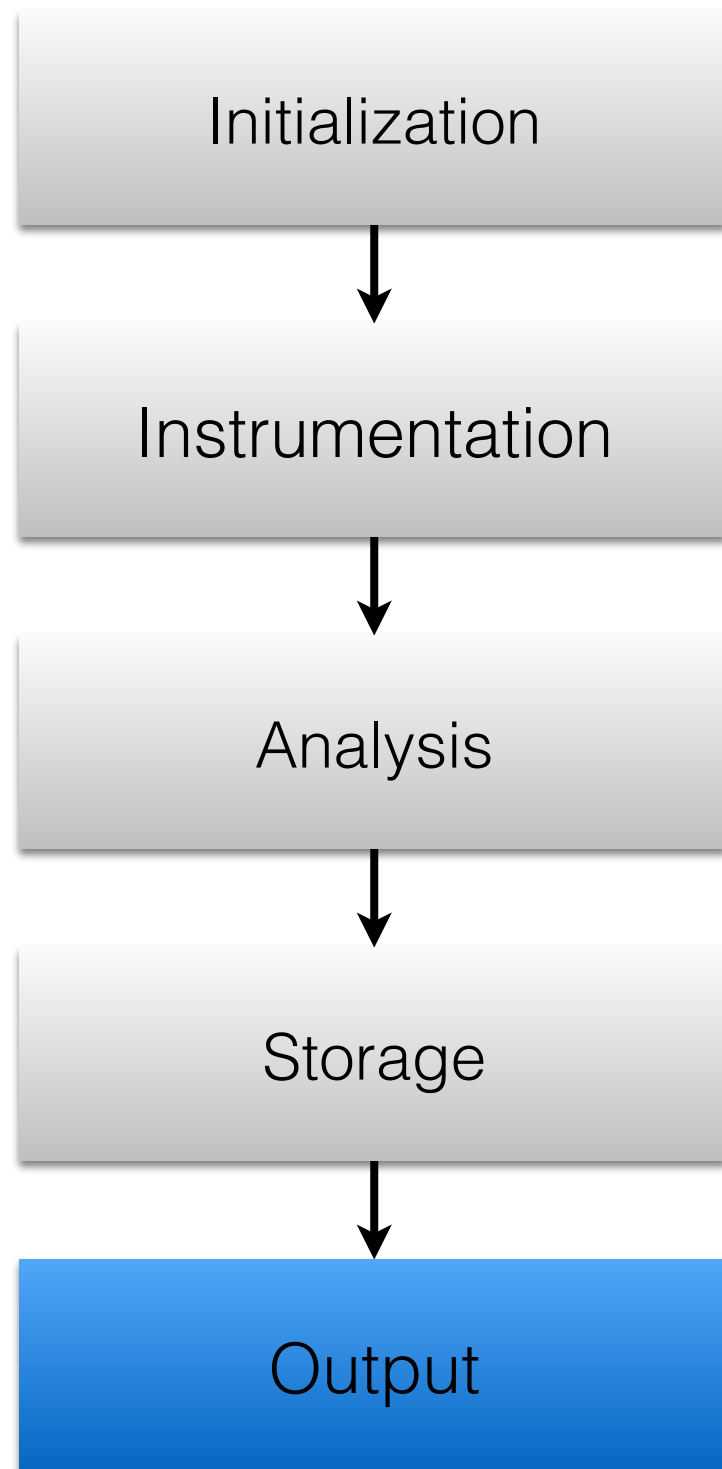
- Analysis code inside the inserted call
  - ♦ User define
  - ♦ Find some “**relation**” between instructions
    - \* redundancy, reuse, ...
- Some common used routines
  - ♦ Get the data structure for current thread
  - ♦ Get the CCT handler of current ins
  - ♦ Instruction pair (**related**)
    - \* <CCT handler1, CCT handler2>

# Work Flow to Prepare DrCCTProf Clients



- Shadow memory
  - ♦ Information corresponding to each memory address
- Map<uint64\_t, uint64\_t>
  - ♦ Key: build from two 32bits CCT handler
  - ♦ Value: frequency that the “**relation**” between the two instructions occur

# Work Flow to Prepare DrCCTProf Clients



- Merge Items in Map
  - ♦ Aggregate metrics that associated with the same source line
- Sort the Map info
  - ♦ Based on the frequency
- Log “**relation**” to file
  - ♦ **drcctlib\_print\_full\_cct**(handler)
  - ♦ Use **dr\_mutex\_lock**(void\* lock) to manipulate one log file for multithreading

# DrCCTProf Tutorial Outline

software download

<https://github.com/Xuhpclab/DrCCTProf>

- DrCCTProf APIs
- General usage of DrCCTProf framework for fine-grained profilers
- DrCCTProf framework internals

# DrCCTProf Tutorial Outline

software download

<https://github.com/Xuhpclab/DrCCTProf>

- DrCCTProf APIs
- General usage of DrCCTProf framework for fine-grained profilers
- DrCCTProf framework internals

# Why Maintain Call Path History?

- To associate context of one event of interest with the context of another event
  - ♦ No priory knowledge of whether the current context will be needed in the future

## Example

- Correctness tools
  - ♦ Data race detection
  - ♦ Taint analysis
  - ♦ Array out of bound detection
- Performance analysis tools
  - ♦ Reuse-distance analysis
  - ♦ Cache simulation
  - ♦ False sharing detection
  - ♦ Redundancy detection (e.g. dead writes)
- Other tools
  - ♦ Debugging, testing, resiliency, replay, etc.

# Scale of Call Paths

	Description	Original program running for 10 minutes
Debuggers	On each break point	$< 10^3$
Performance analysis tools	On each sample (1 sample/ms)	$6 \times 10^5$
Fine-grained instrumentation tools	On each instruction (2GHz CPU)	$1.2 \times 10^{12}$

# State-of-the-art in Collecting Ubiquitous Call Paths



**“It will slow down execution by a factor of several thousand compared to native execution -- I'd guess -- so you'll wind up with something that is unusably slow on anything except the smallest problems.”**

*Dyn*  
*inst*

**“If you tried to invoke `Thread::getCallStack` on every memory access there would be very serious performance problems ... your program would probably never reach main.”**



**No support for collecting calling contexts**



# State-of-the-art in Collecting Ubiquitous Call Paths



“It will slow down execution by a factor of several thousand compared to native execution -- I'd guess -- so you'll wind up with something that is unusably slow on anything except the smallest problems.”

*Dyn*  
*inst*

“If you tried to invoke `Thread::getCallStack` on every memory access there would be very serious performance problems ... your program would probably never reach main.”



No support for collecting calling contexts

We built one ourselves—**DrCCTProf**

# Challenges in Ubiquitous Call Path Collection

1. Overhead (space)
2. Overhead (time)
3. Overhead (parallel scaling)

# Store History of Contexts Compactly

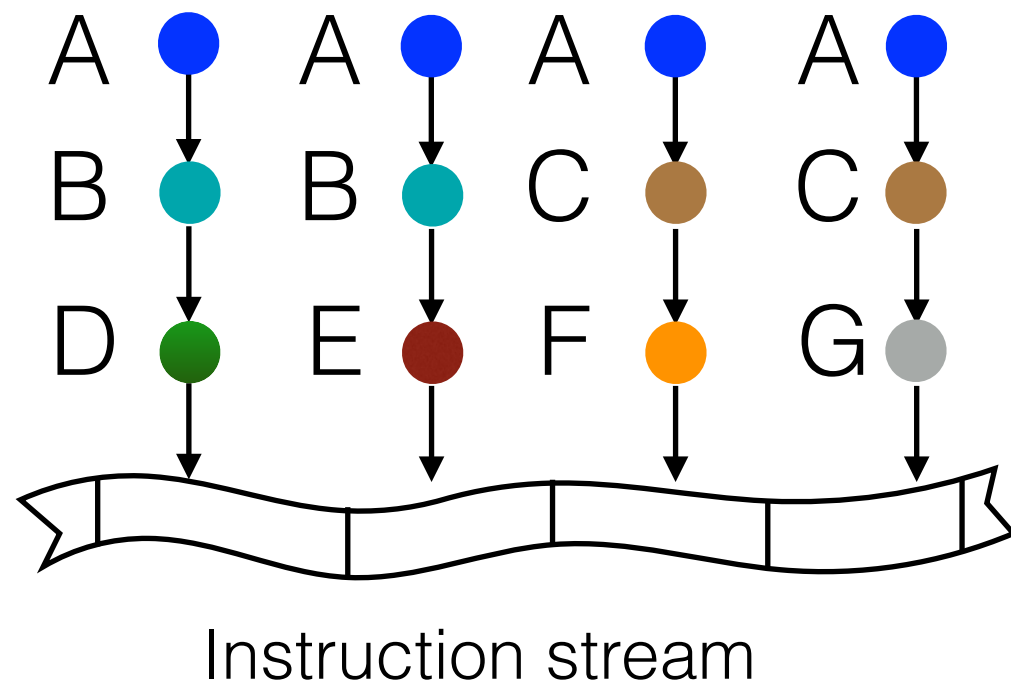
Problem:

Deluge of call paths

# Store History of Contexts Compactly

Problem:

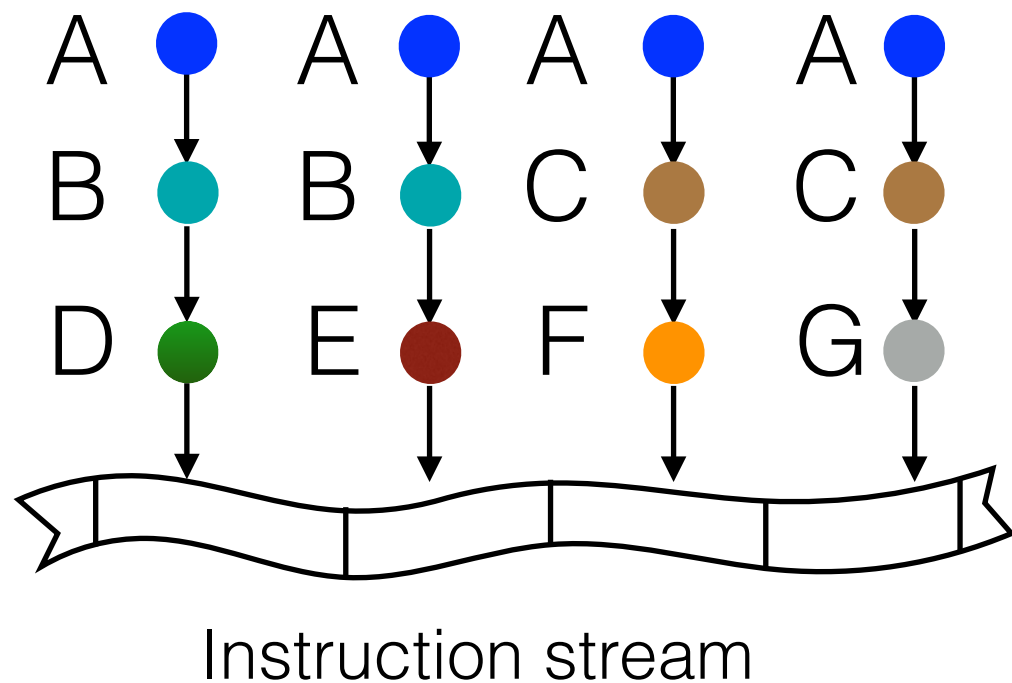
Deluge of call paths



# Store History of Contexts Compactly

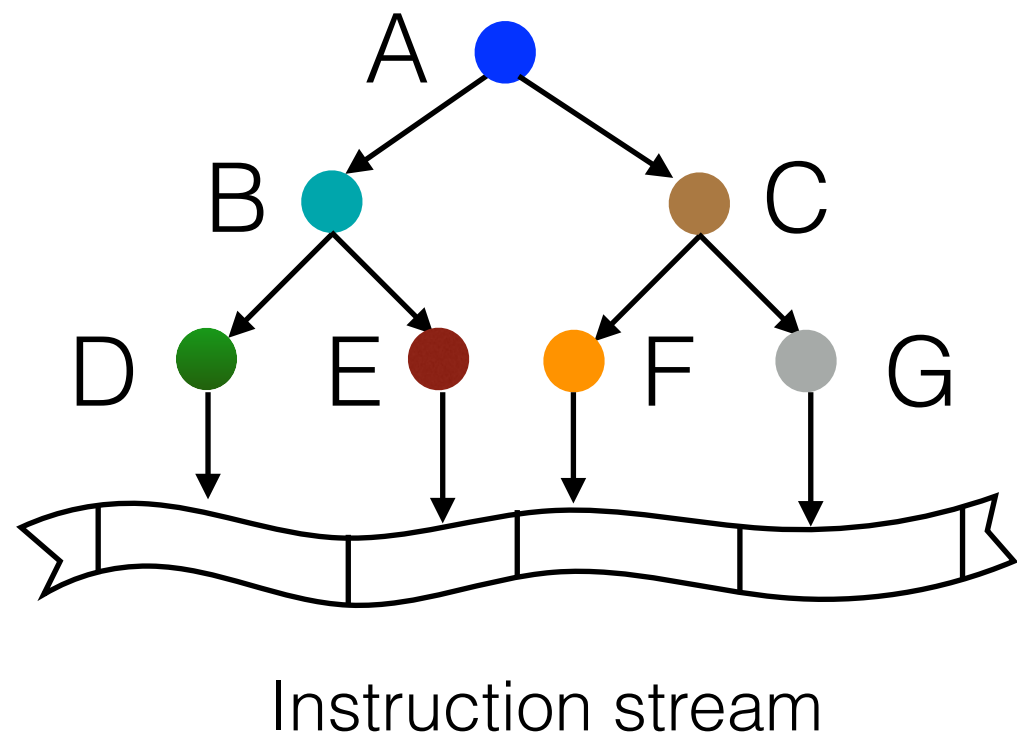
## Problem:

Deluge of call paths



## Solution

- Call paths share common prefix
- Store call paths as a calling context tree (CCT)
- One CCT per thread



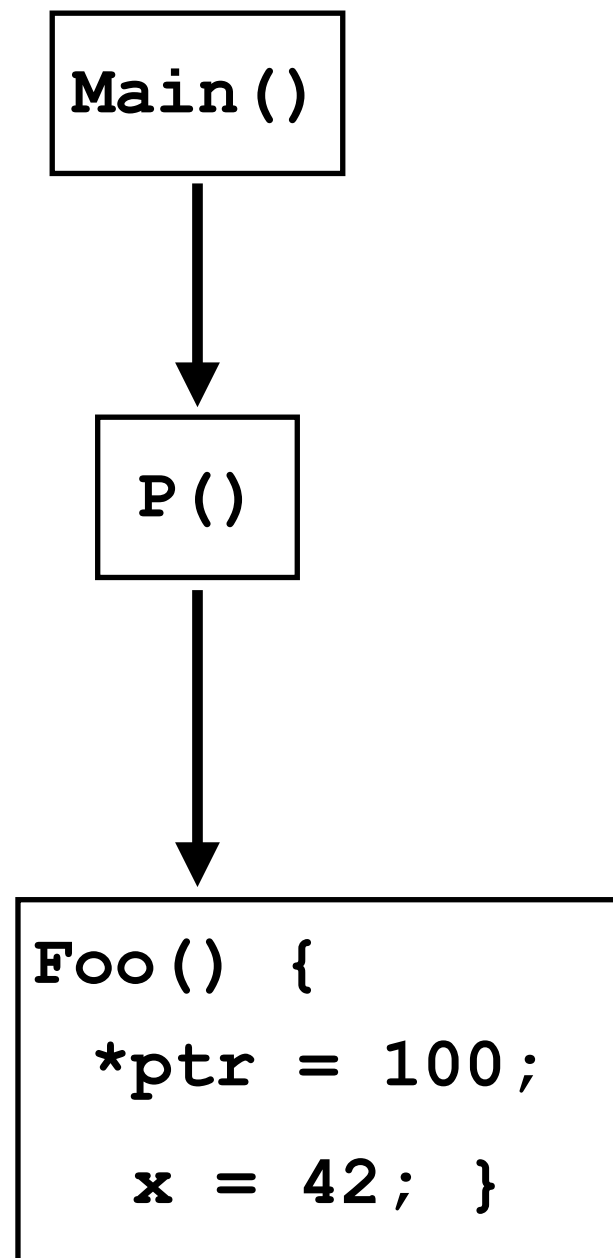
# Shadow Stack to Avoid Unwinding Overhead

Problem:

Unwinding overhead

Solution:

Reverse the process. Eagerly build a replica/shadow stack on-the-fly.



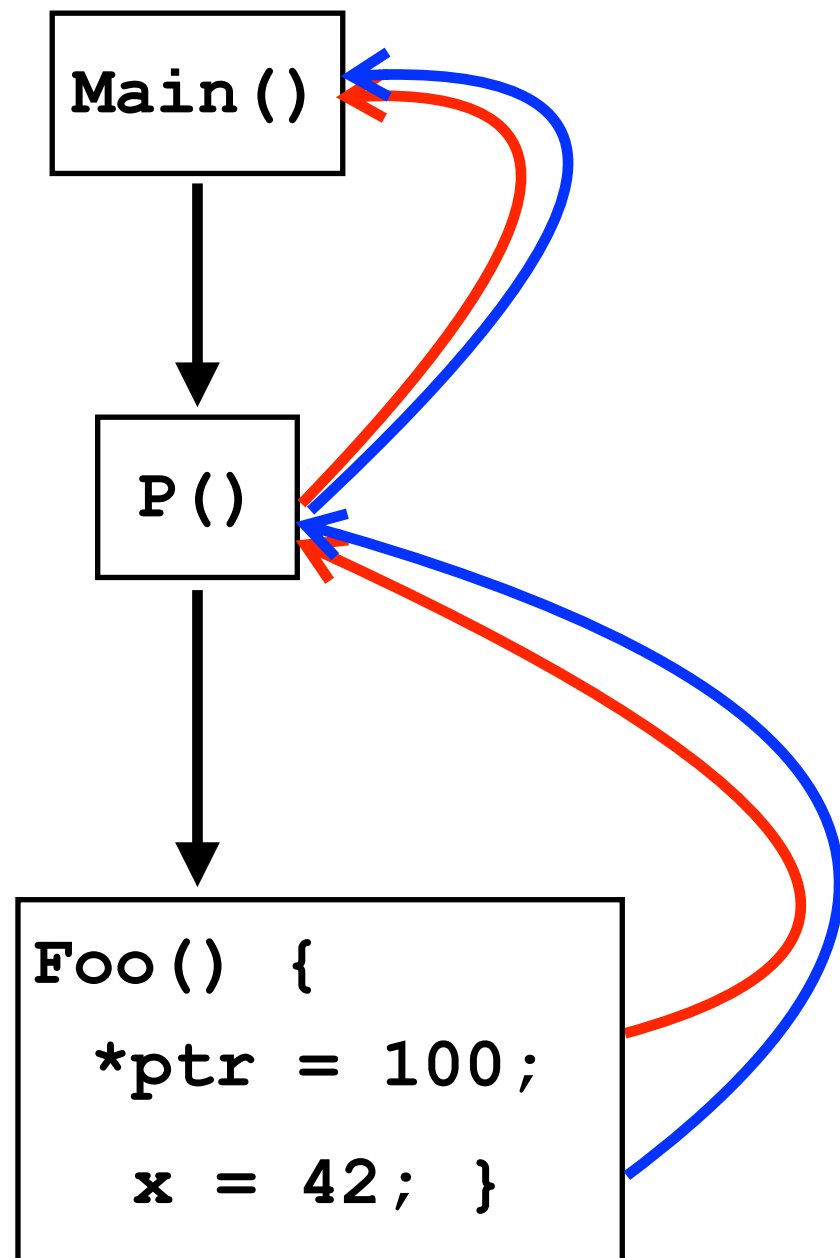
# Shadow Stack to Avoid Unwinding Overhead

Problem:

Unwinding overhead

Solution:

Reverse the process. Eagerly build a replica/shadow stack on-the-fly.



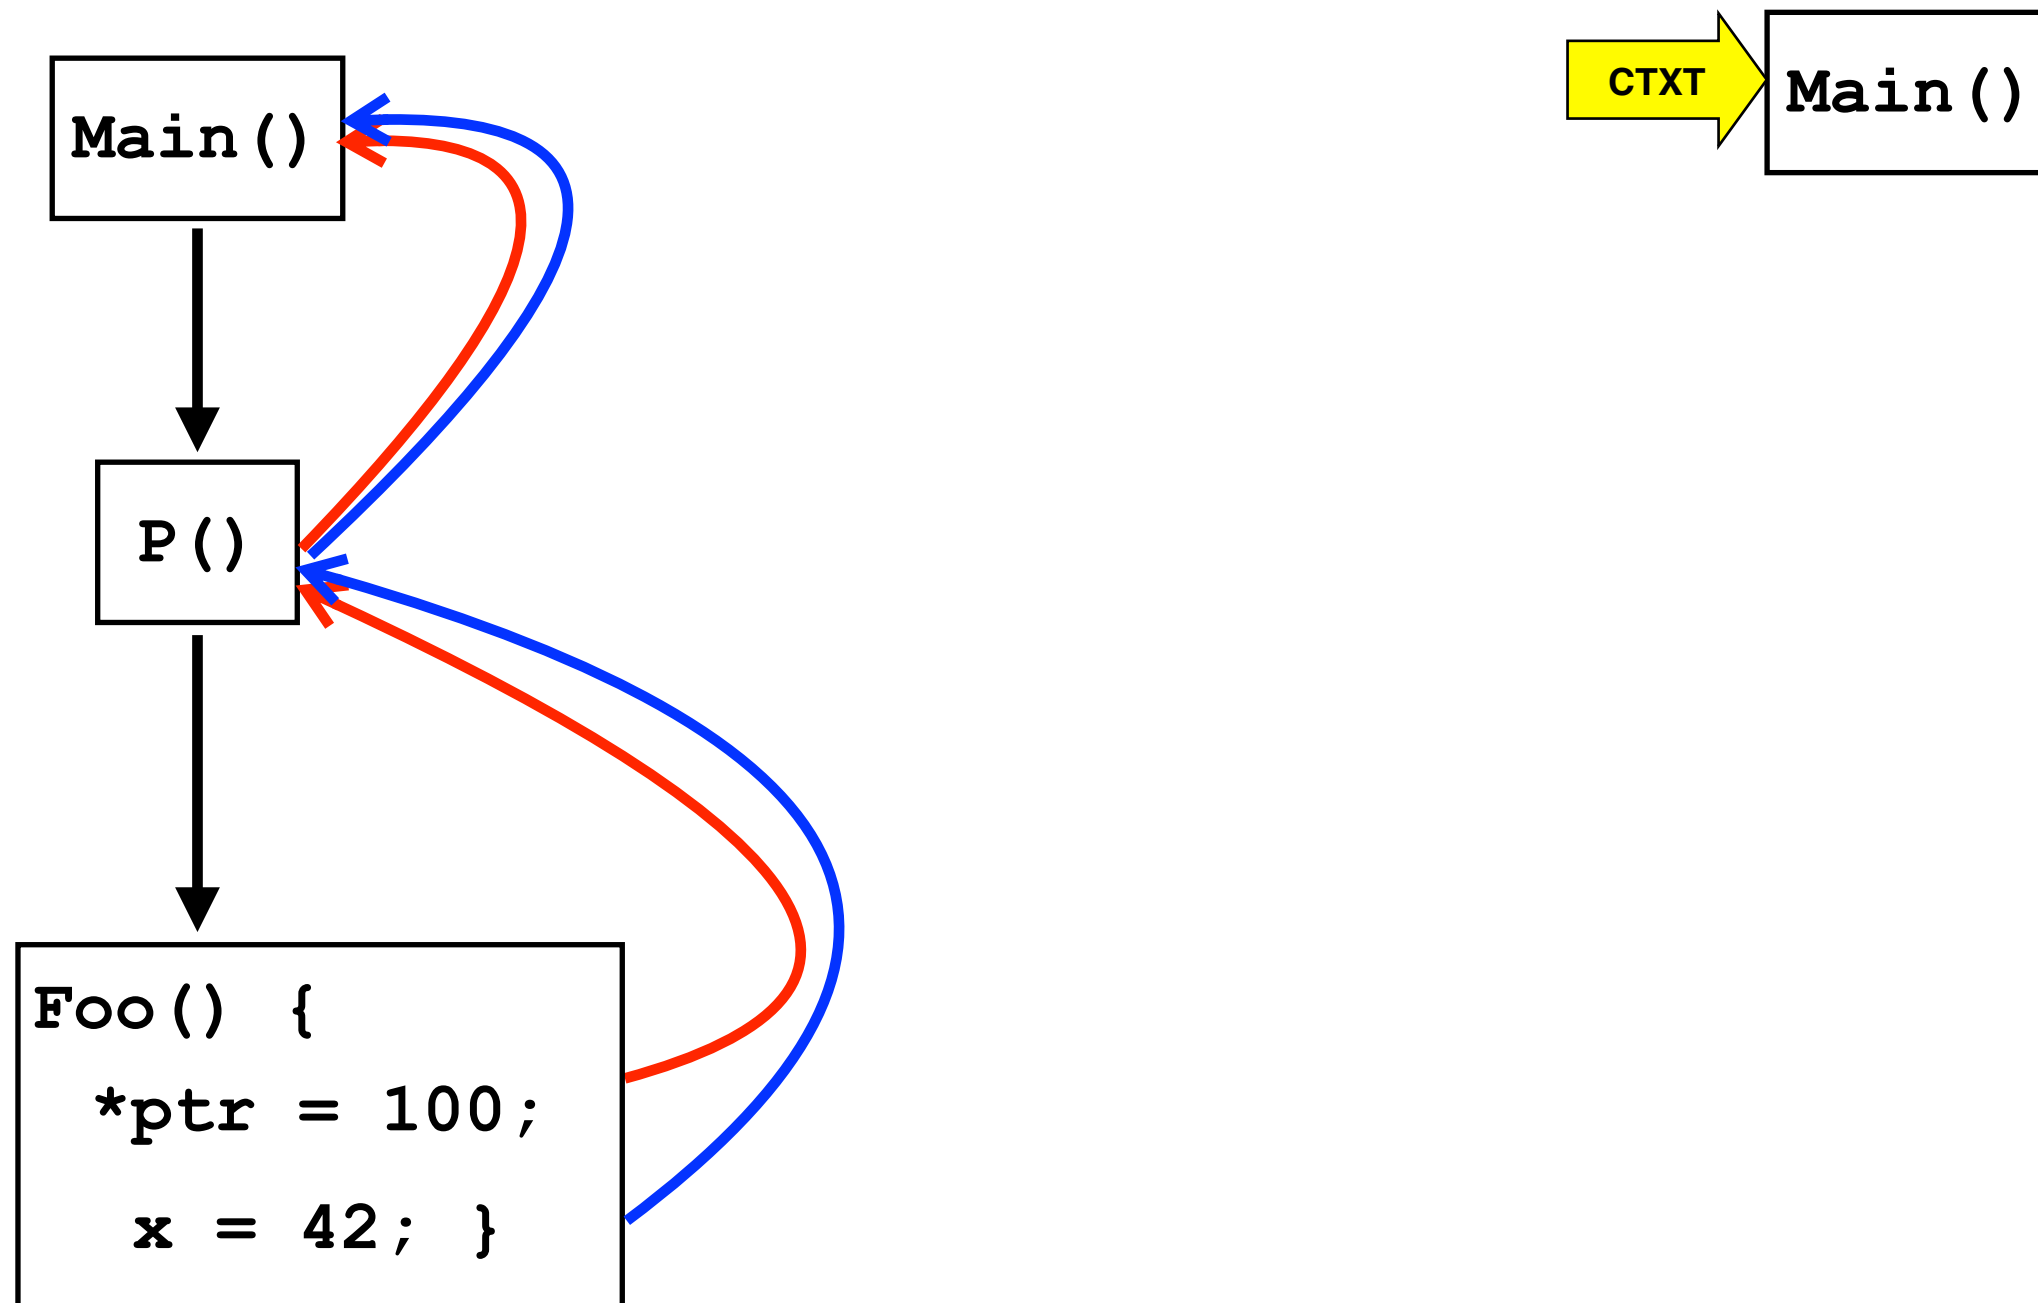
# Shadow Stack to Avoid Unwinding Overhead

Problem:

Unwinding overhead

Solution:

Reverse the process. Eagerly build a replica/shadow stack on-the-fly.





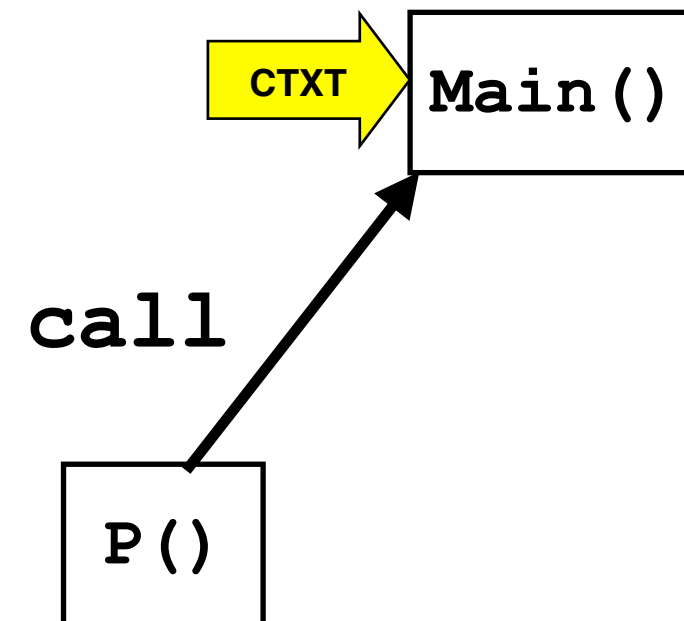
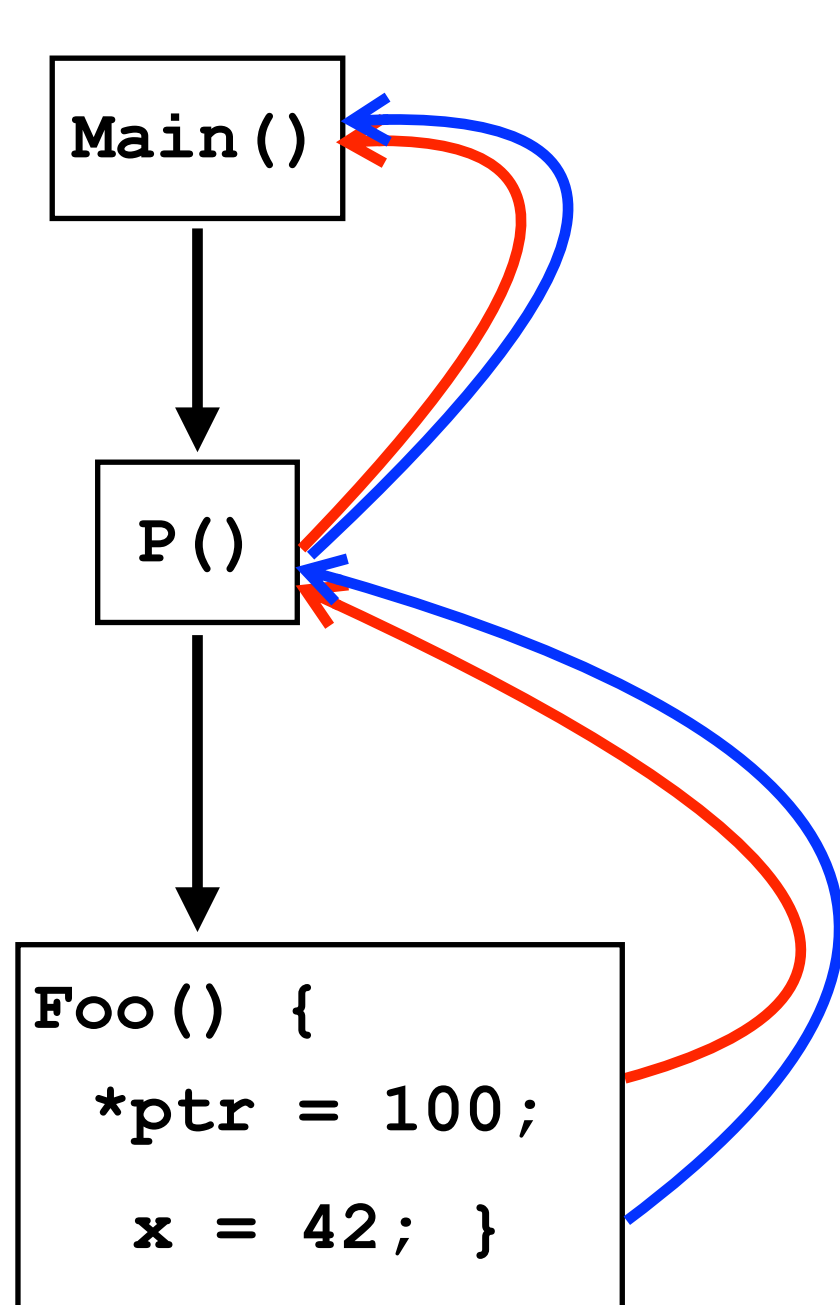
# Shadow Stack to Avoid Unwinding Overhead

Problem:

Unwinding overhead

Solution:

Reverse the process. Eagerly build a replica/shadow stack on-the-fly.



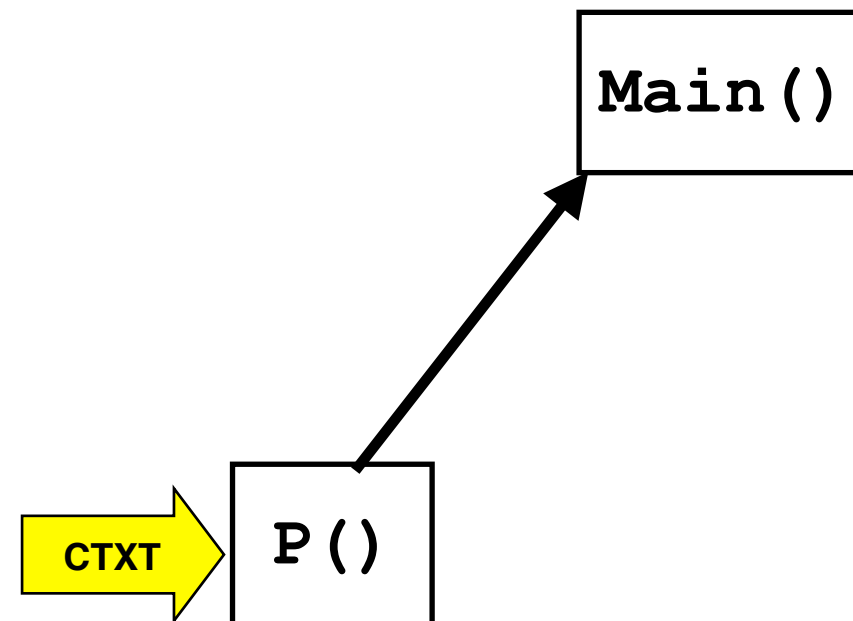
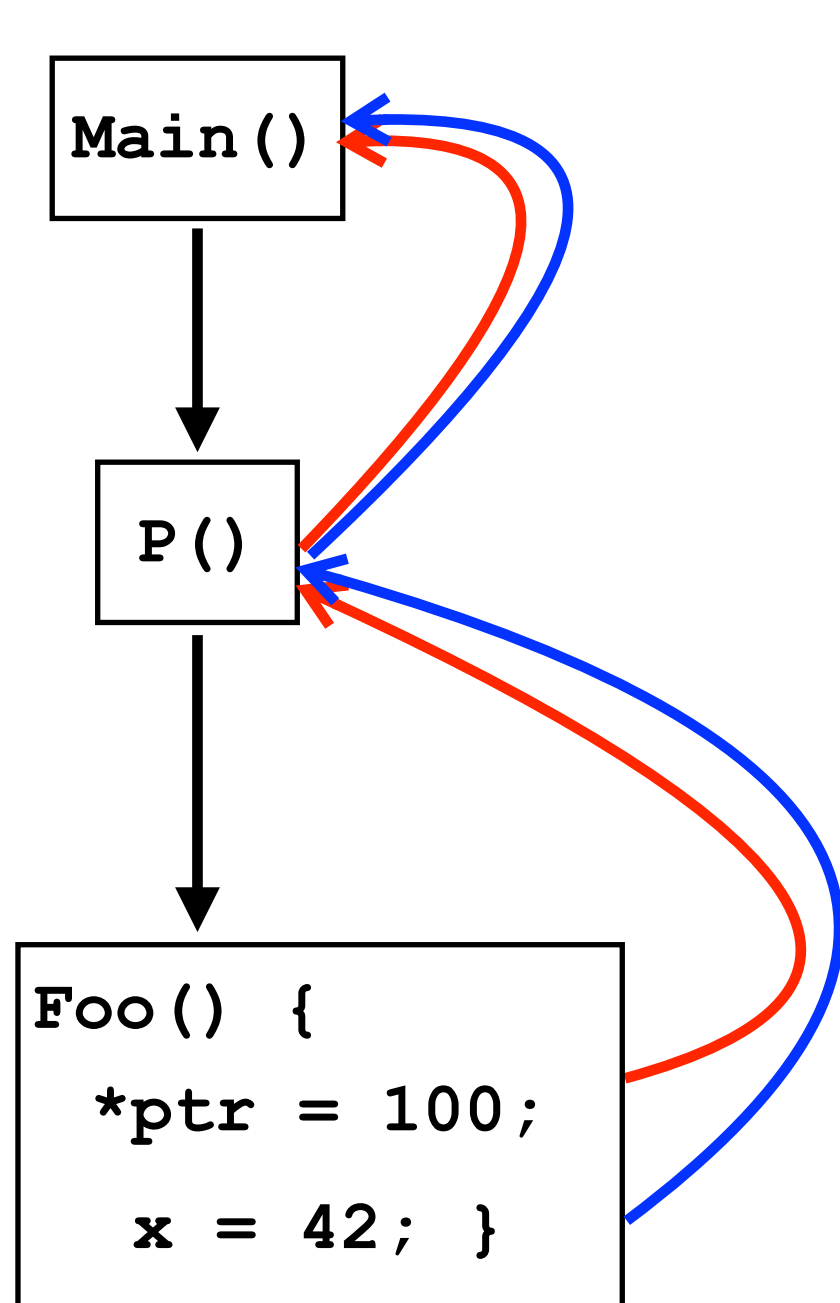
# Shadow Stack to Avoid Unwinding Overhead

Problem:

Unwinding overhead

Solution:

Reverse the process. Eagerly build a replica/shadow stack on-the-fly.



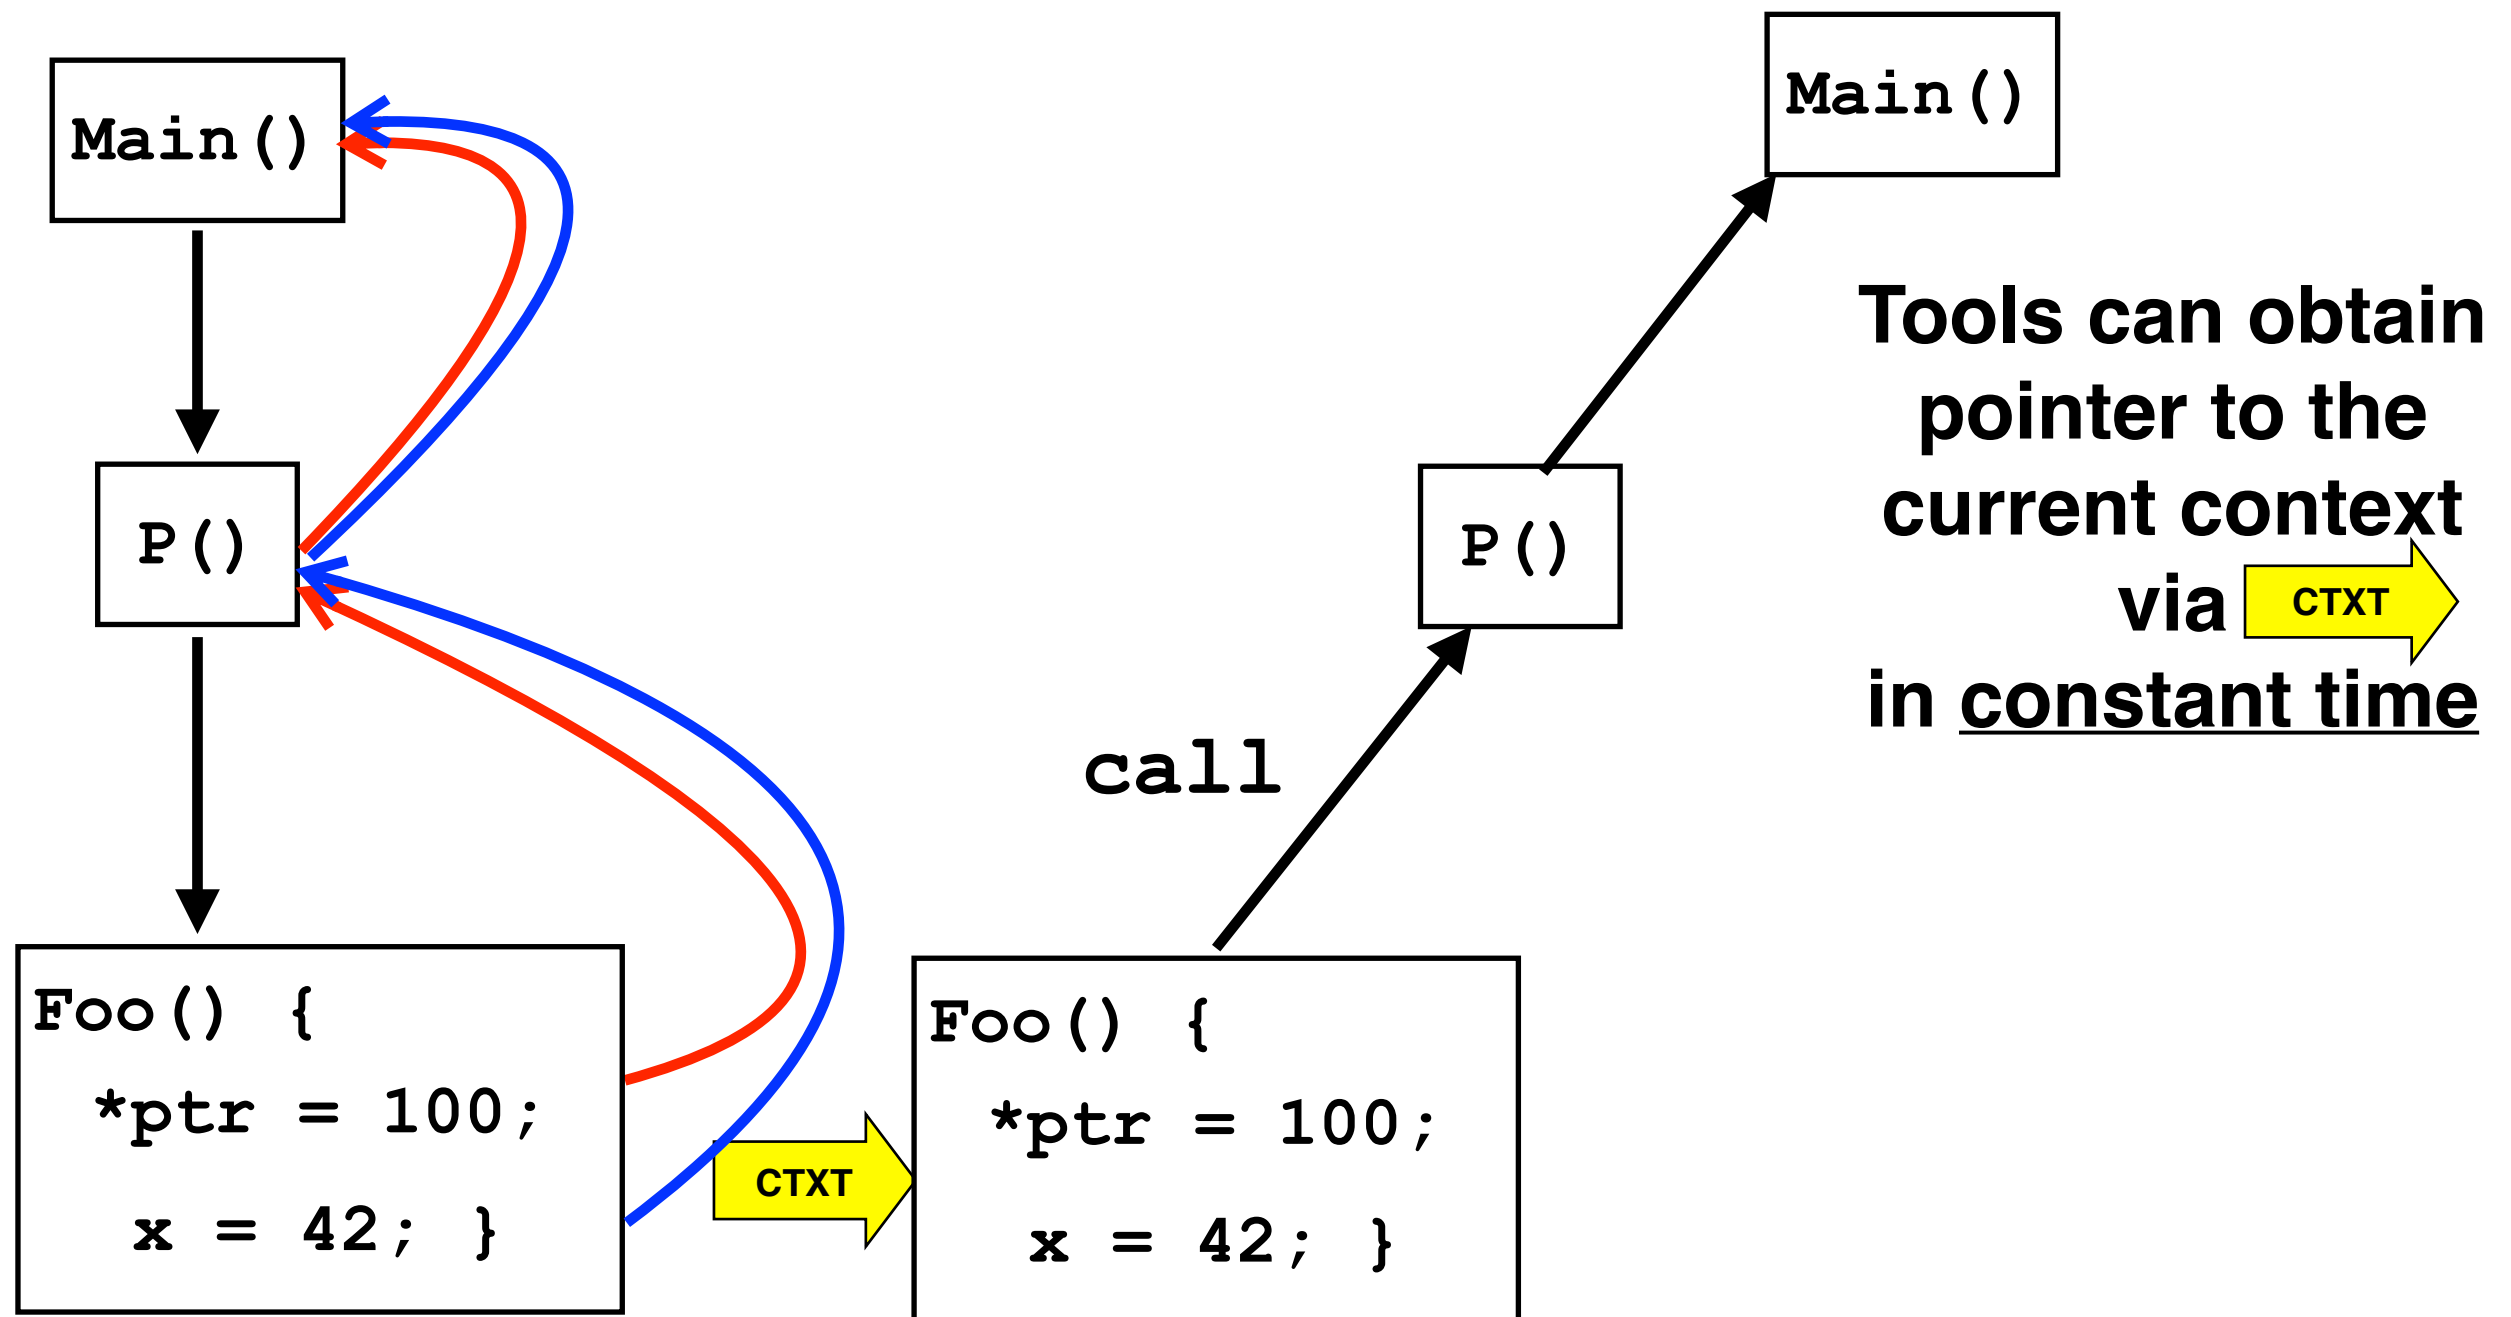
# Shadow Stack to Avoid Unwinding Overhead

Problem:

Unwinding overhead

Solution:

Reverse the process. Eagerly build a replica/shadow stack on-the-fly.



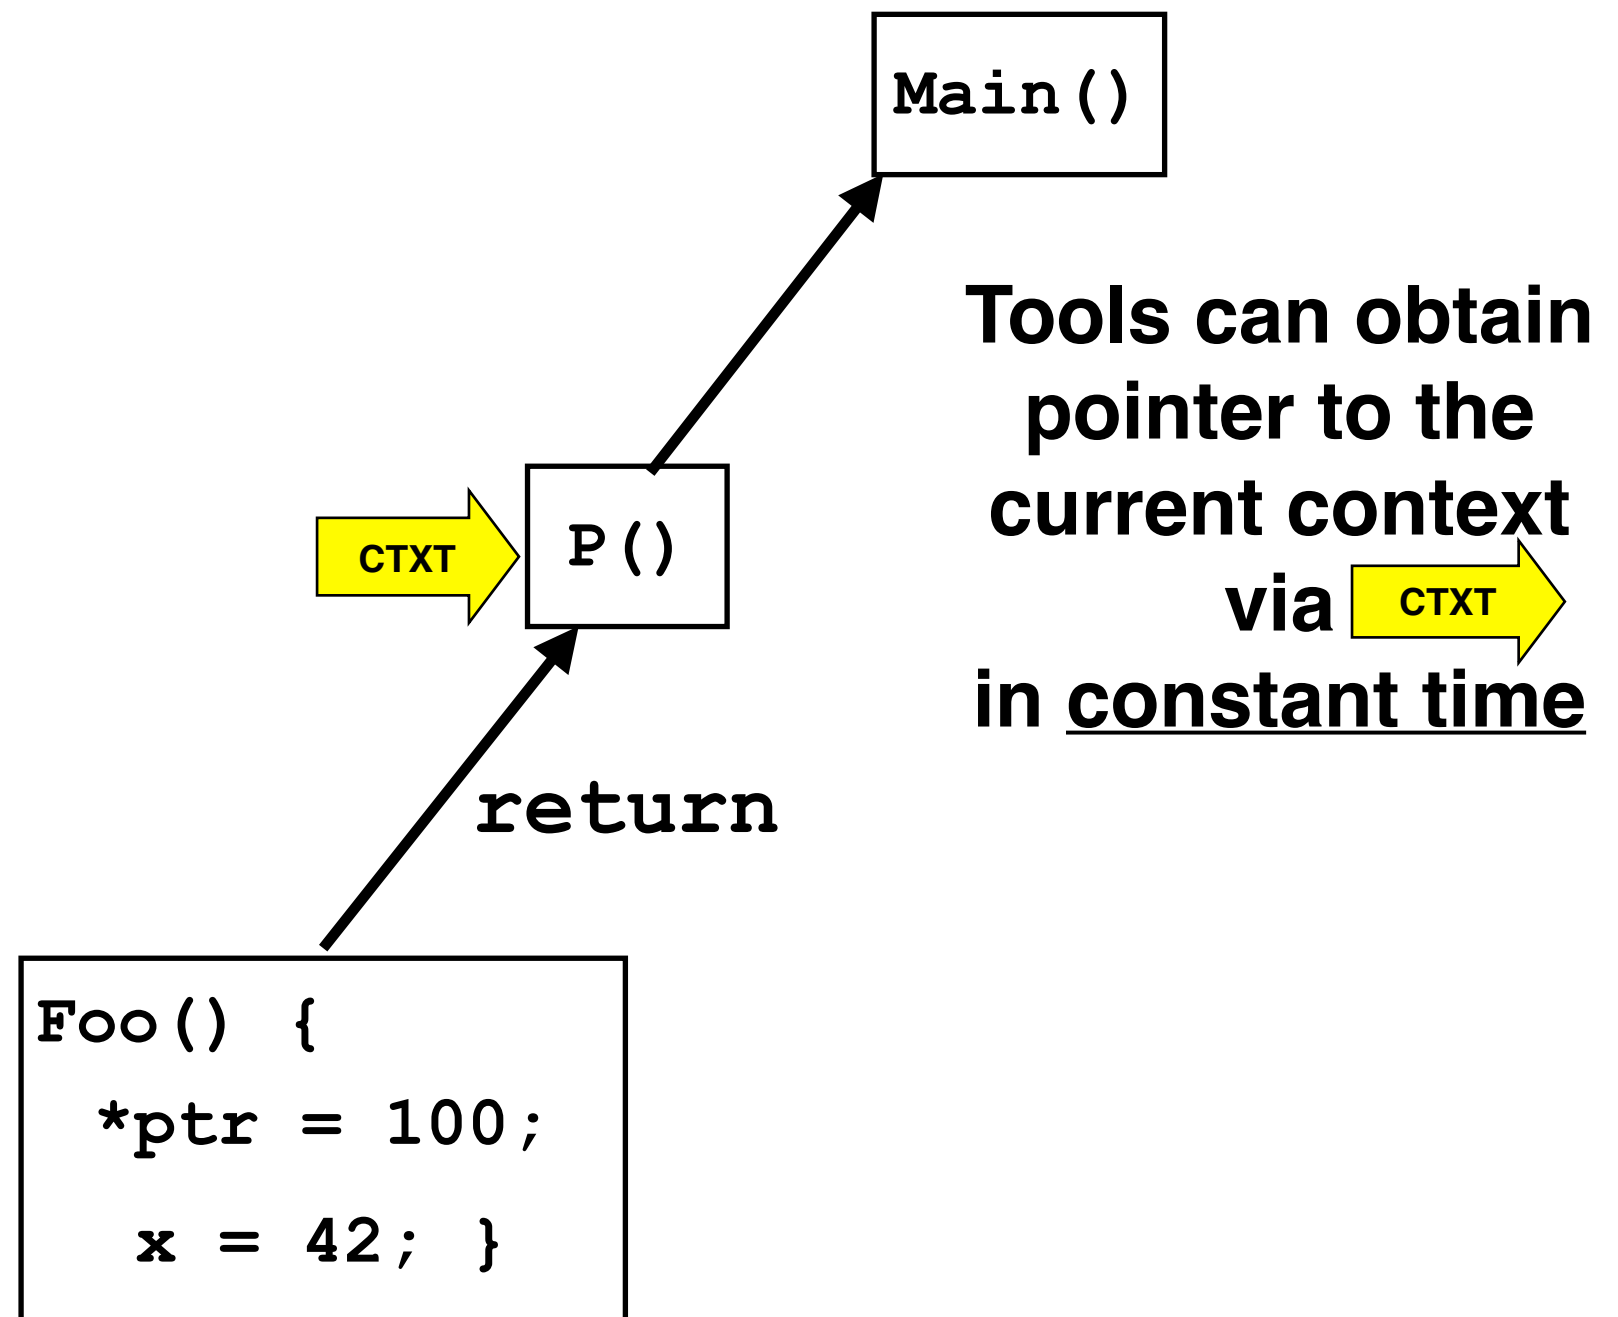
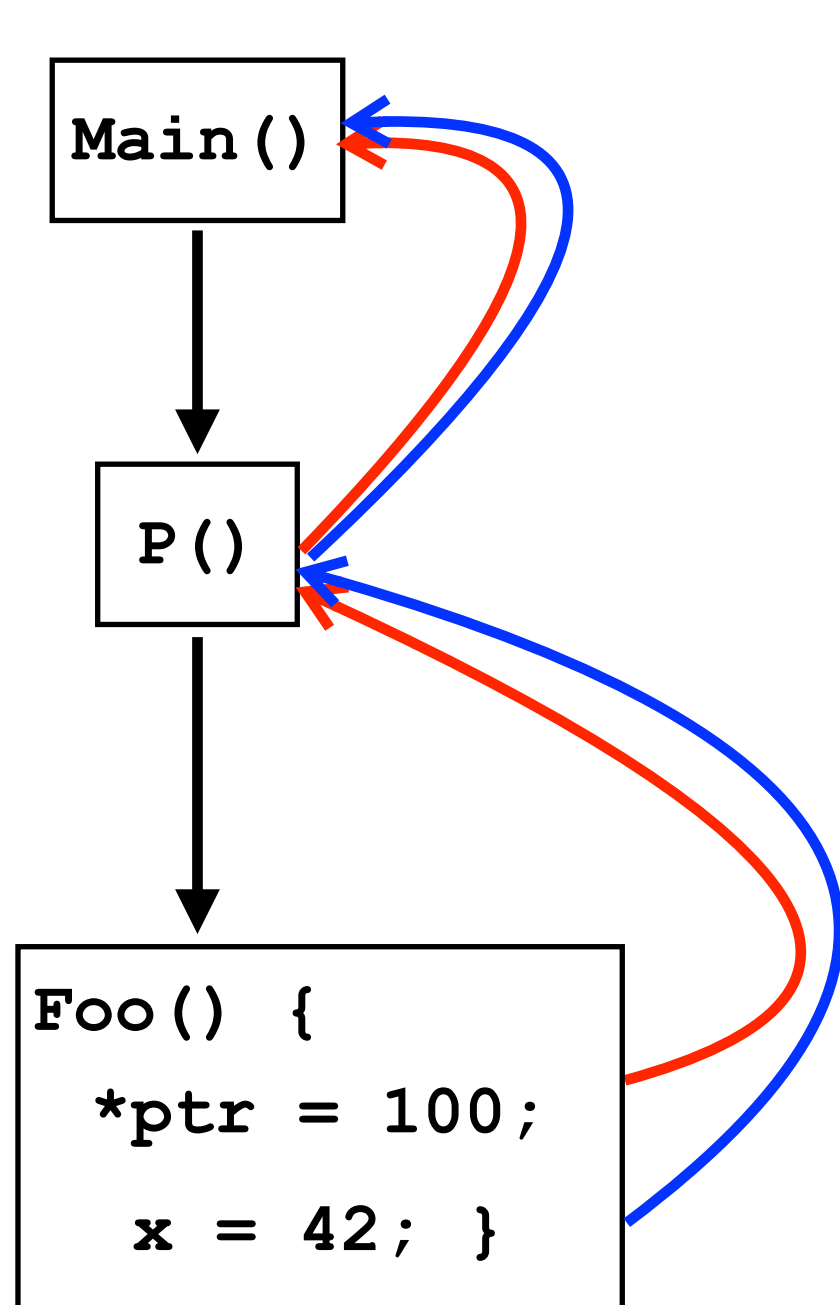
# Shadow Stack to Avoid Unwinding Overhead

Problem:

Unwinding overhead

Solution:

Reverse the process. Eagerly build a replica/shadow stack on-the-fly.



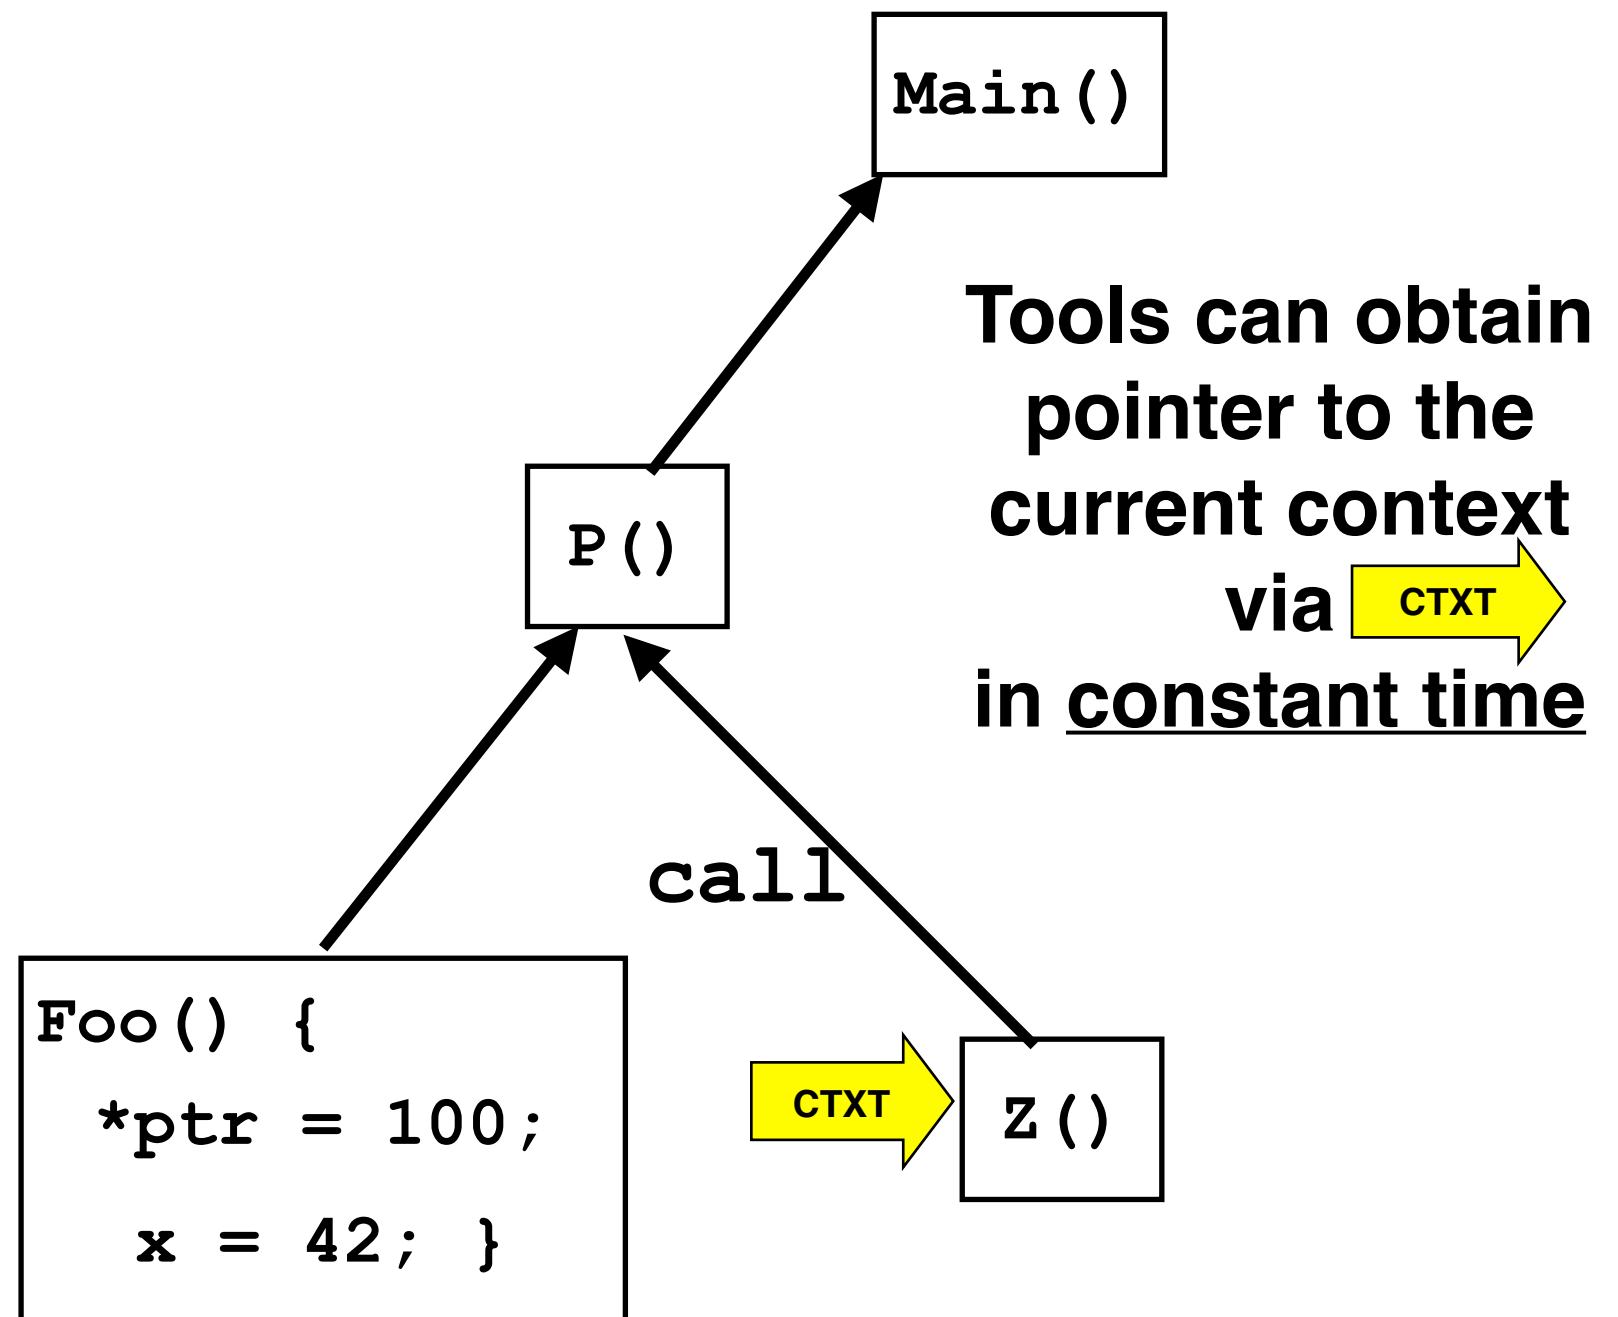
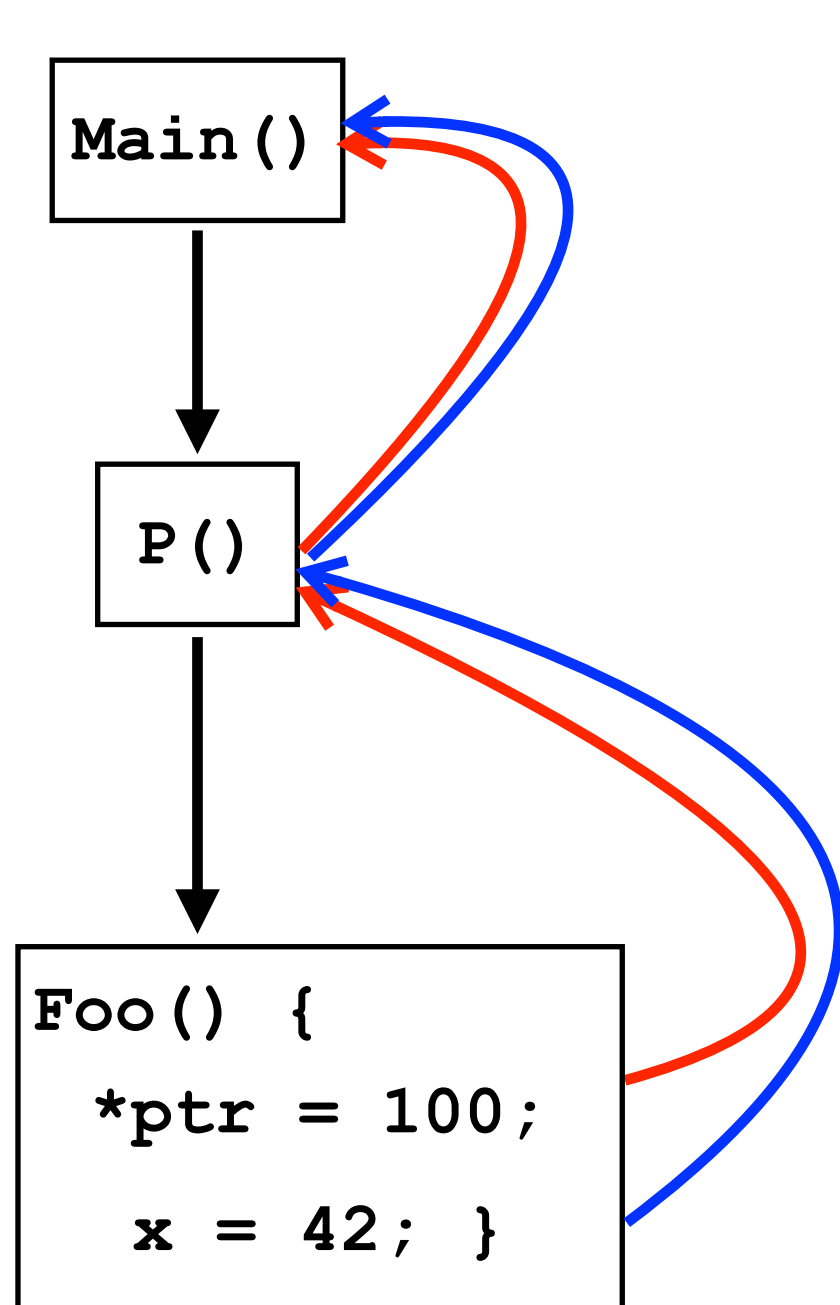
# Shadow Stack to Avoid Unwinding Overhead

Problem:

Unwinding overhead

Solution:

Reverse the process. Eagerly build a replica/shadow stack on-the-fly.



# Complexities of Context Management

- Problem #1: Call/Return instructions are not simple to spot
  - ♦ Static disassembly is imperfect
    - \* Data embedded in instruction stream
    - \* Missing or incorrect function boundary information
- Problem #2: CTEXT needs call-path + PC for completeness
  - ♦ Naively using the <Context:PC> tuple makes each handle >64 bits
  - ♦ We wish to keep a handle 32-bit since tools often use two 32-bit (64-bit) handles as a hash-table index
- Solution: basic block instrumentation
  - ♦ DynamoRIO instruments every BB
  - ♦ Call/return instructions discovered in a BB are reliable
  - ♦ DrCCTProf assigns a unique 32-bit id to each instruction appearing in each call path in each trace —> supports 4.2B unique contexts (including PC)

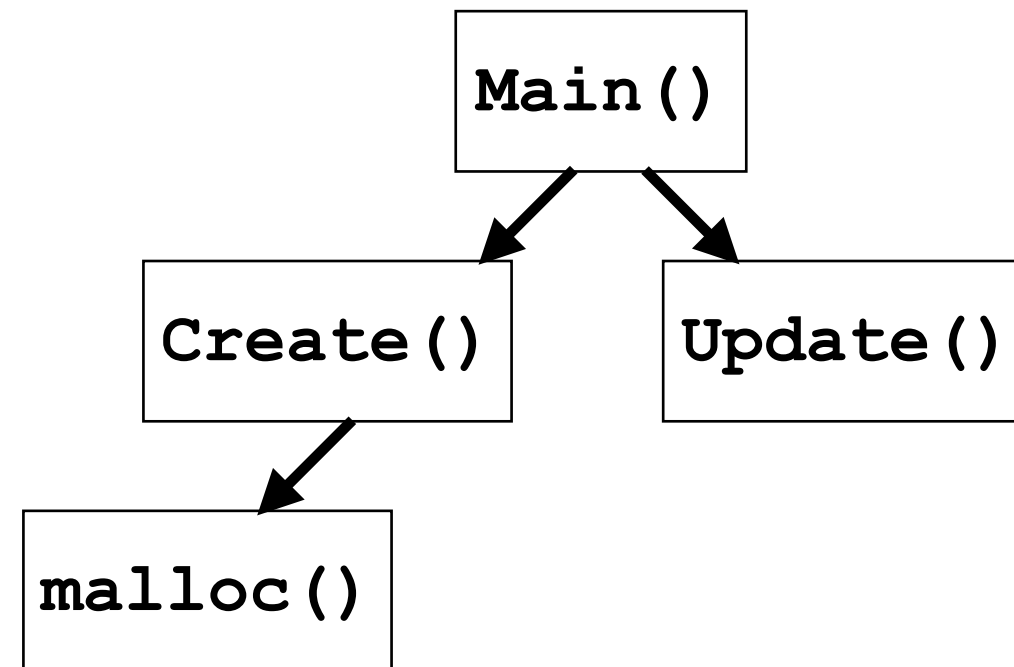
# Associating Address With Data Objects

```
int MyArray[SZ];

int * Create() {
    return malloc(...);
}

void Update(int * ptr) {
    for( ... )
        ptr[i]++;
}

int main() {
    int * p;
    if (...)
        p = Create();
    else
        p = MyArray;
    Update(p);
}
```



- Associate each *data access* to its *data object*
- Data object
  - ♦ Dynamic allocation: Call path of allocation site
  - ♦ Static objects: Variable name

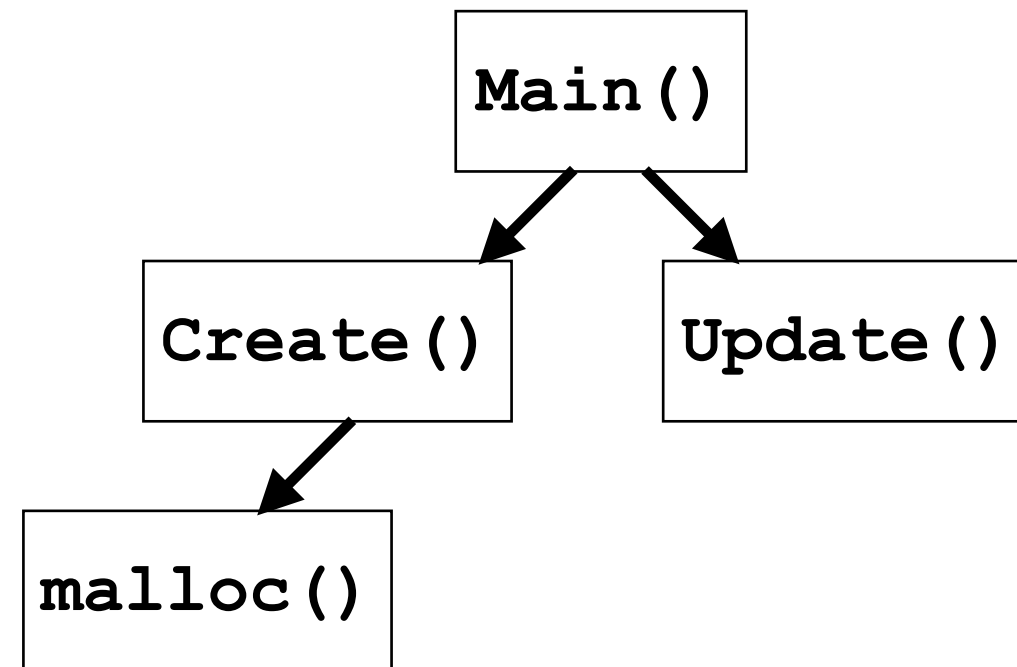
# Associating Address With Data Objects

```
int MyArray[SZ];

int * Create() {
    return malloc(...);
}

void Update(int * ptr) {
    for( ... )
        ptr[i]++;
}

int main() {
    int * p;
    if (...)
        p = Create();
    else
        p = MyArray;
    Update(p);
}
```



- Associate each *data access* to its *data object*
- Data object
  - ♦ Dynamic allocation: Call path of allocation site
  - ♦ Static objects: Variable name



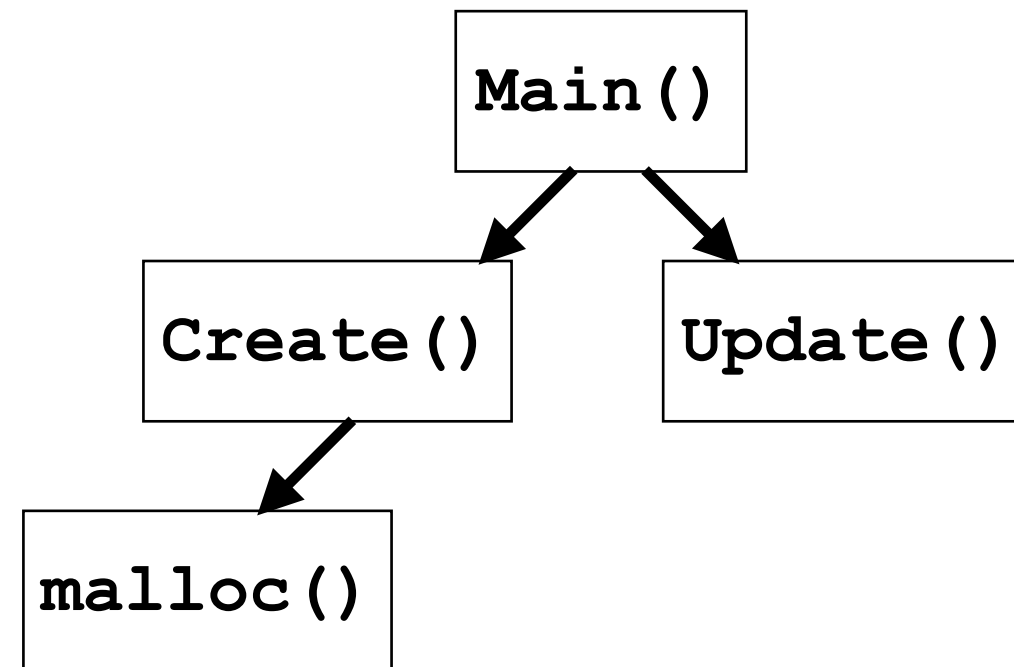
# Associating Address With Data Objects

```
int MyArray[SZ];

int * Create() {
    return malloc(...);
}

void Update(int * ptr) {
    for( ... )
        ptr[i]++;
}

int main() {
    int * p;
    if (...)
        p = Create();
    else
        p = MyArray;
    Update(p);
}
```



- Associate each *data access* to its *data object*
- Data object
  - ♦ Dynamic allocation: Call path of allocation site
  - ♦ Static objects: Variable name

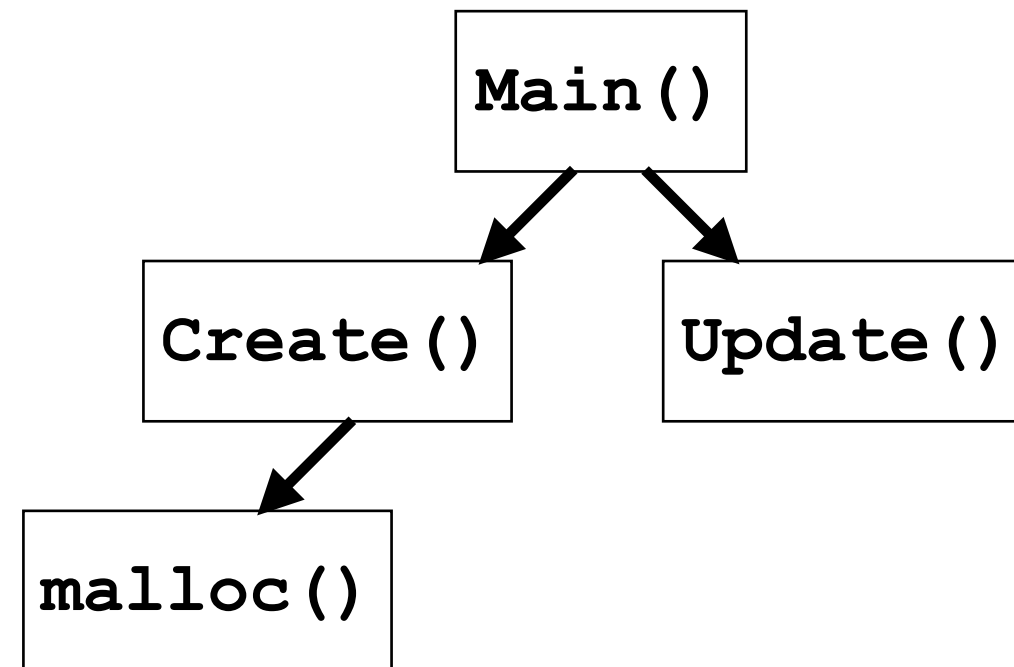
# Associating Address With Data Objects

```
int MyArray[SZ];

int * Create() {
    return malloc(...);
}

void Update(int * ptr) {
    for( ... )
        ptr[i]++;
}

int main() {
    int * p;
    if (...)
        p = Create();
    else
        p = MyArray;
    Update(p);
}
```



- Associate each *data access* to its *data object*
- Data object
  - ♦ Dynamic allocation: Call path of allocation site
  - ♦ Static objects: Variable name

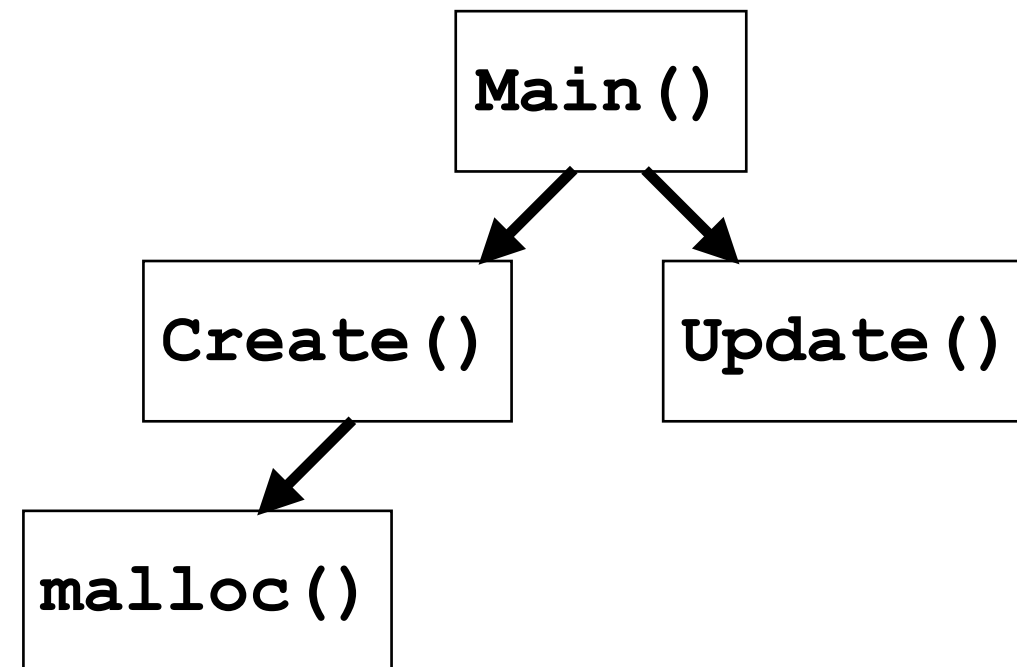
# Associating Address With Data Objects

```
int MyArray[SZ];

int * Create() {
    return malloc(...);
}

void Update(int * ptr) {
    for( ... )
        ptr[i]++; ←
}

int main() {
    int * p;
    if (...)
        p = Create();
    else
        p = MyArray;
    Update(p);
}
```



- Associate each *data access* to its *data object*
- Data object
  - ♦ Dynamic allocation: Call path of allocation site
  - ♦ Static objects: Variable name

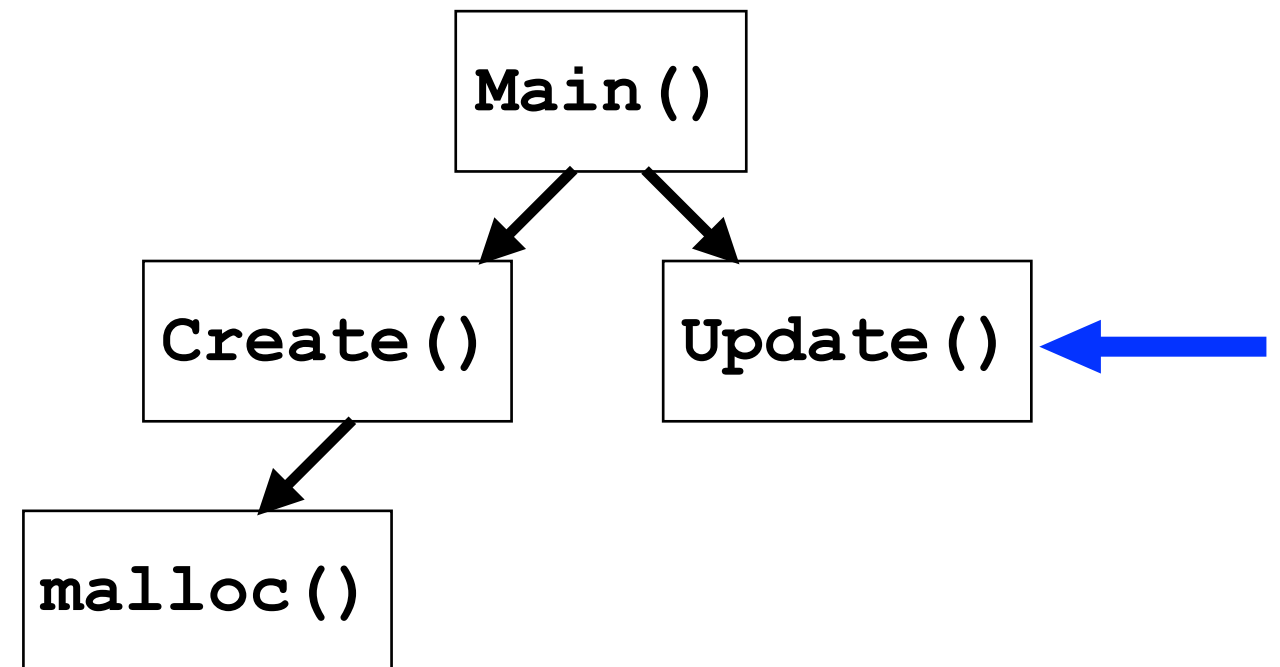
# Associating Address With Data Objects

```
int MyArray[SZ];

int * Create() {
    return malloc(...);
}

void Update(int * ptr) {
    for( ... )
        ptr[i]++;
}

int main() {
    int * p;
    if (...)
        p = Create();
    else
        p = MyArray;
    Update(p);
}
```



- Associate each *data access* to its *data object*
- Data object
  - ♦ Dynamic allocation: Call path of allocation site
  - ♦ Static objects: Variable name

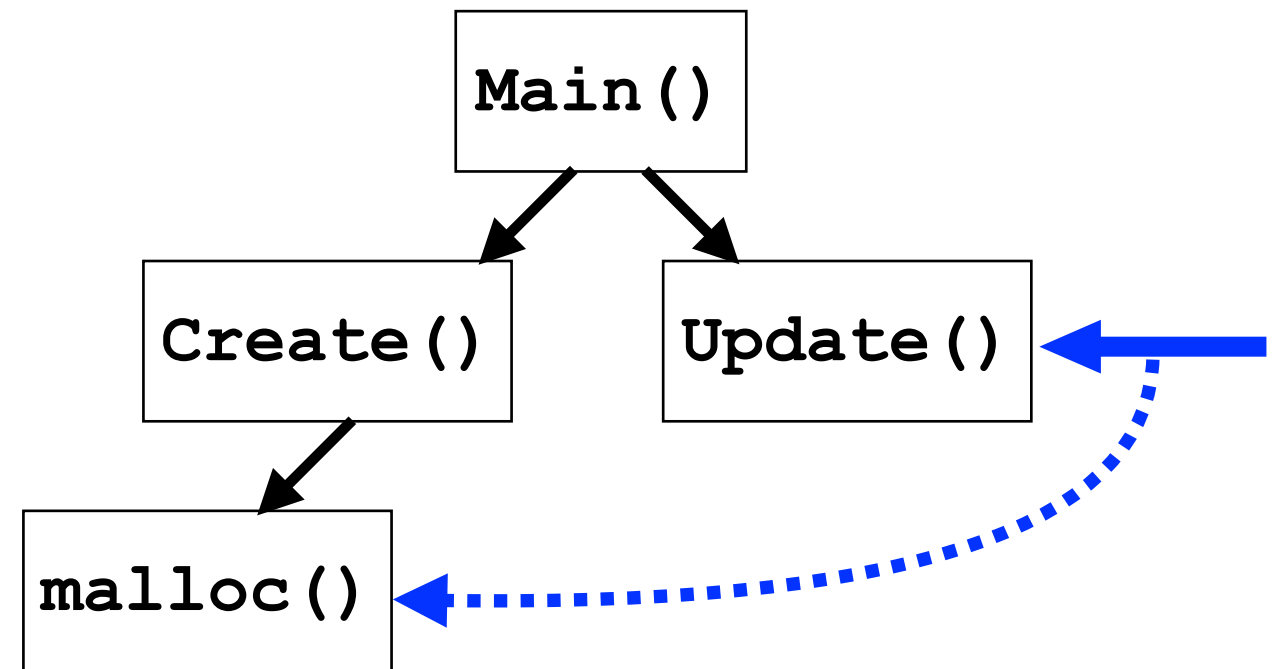
# Associating Address With Data Objects

```
int MyArray[SZ];

int * Create() {
    return malloc(...);
}

void Update(int * ptr) {
    for( ... )
        ptr[i]++;
}

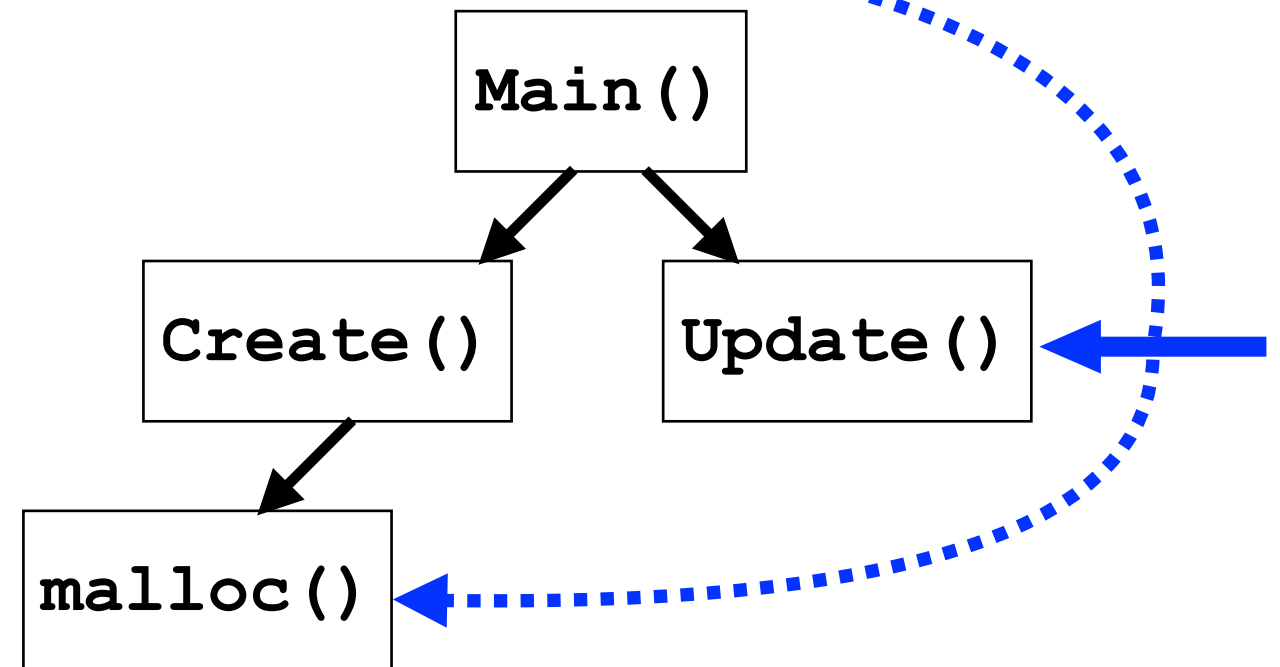
int main() {
    int * p;
    if (...)
        p = Create();
    else
        p = MyArray;
    Update(p);
}
```



- Associate each *data access* to its *data object*
- Data object
  - ✦ Dynamic allocation: Call path of allocation site
  - ✦ Static objects: Variable name

# Associating Address With Data Objects

```
int MyArray[SZ];  
  
int * Create() {  
    return malloc(...);  
}  
  
void Update(int * ptr) {  
    for( ... )  
        ptr[i]++;  
}  
  
int main() {  
    int * p;  
    if (...)  
        p = Create();  
    else  
        p = MyArray;  
    Update(p);  
}
```



- Associate each *data access* to its *data object*
- Data object
  - ♦ Dynamic allocation: Call path of allocation site
  - ♦ Static objects: Variable name

# How to Associate Address with Data Objects

- Static objects
  - ♦ Record all `<AddressRange, VariableName>` tuples in a map
- Dynamic allocations
  - ♦ Instrument all allocation/free routines
  - ♦ Maintain `<AddressRange, ContextId>` tuples in the map
- At each memory access: search the map for the address
- Problems
  - ♦ Searching the map on each access is expensive
  - ♦ Map needs to be concurrent for threaded programs

# Address→Object Mapping via Shadow Memory

Solution: shadow memory

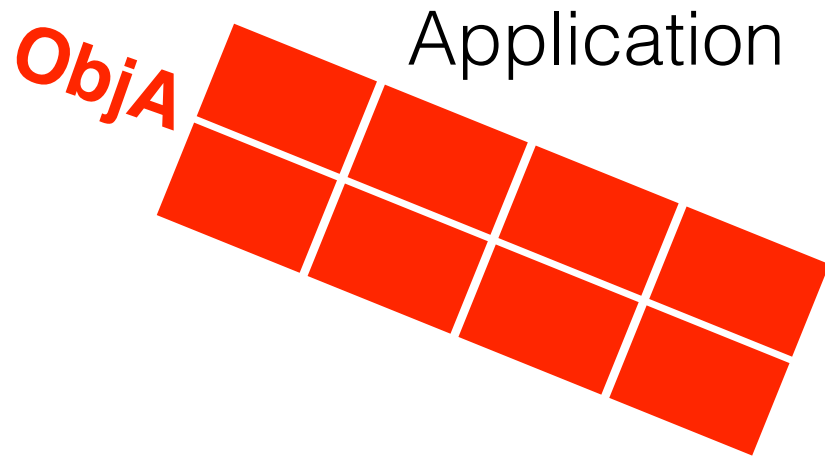
Application

DrCCTProf



# Address→Object Mapping via Shadow Memory

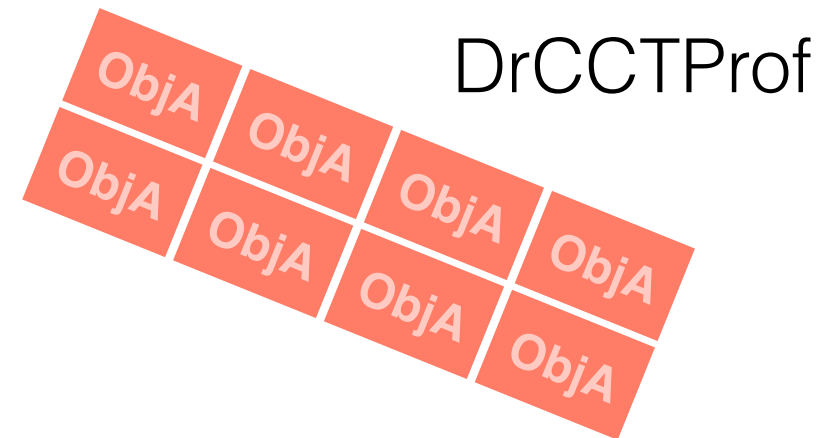
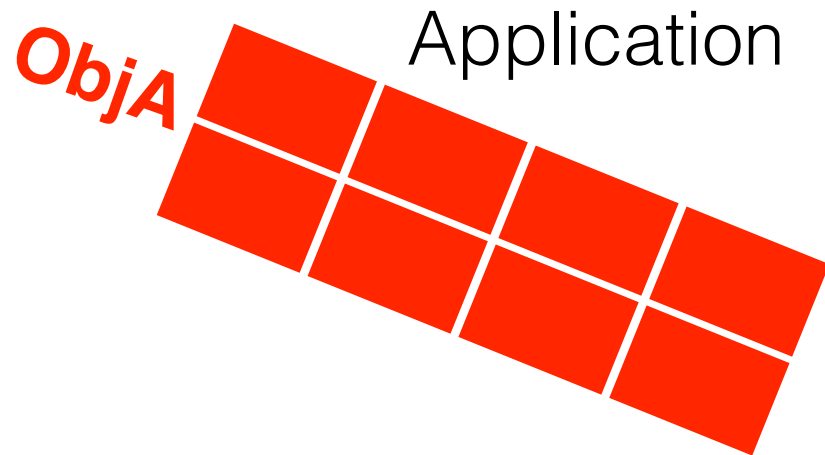
Solution: shadow memory



DrCCTProf

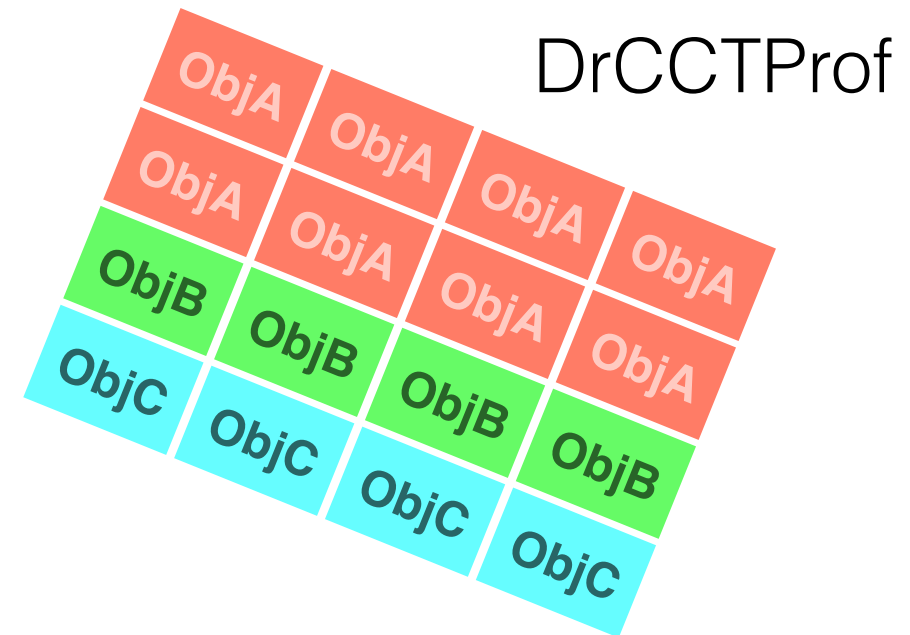
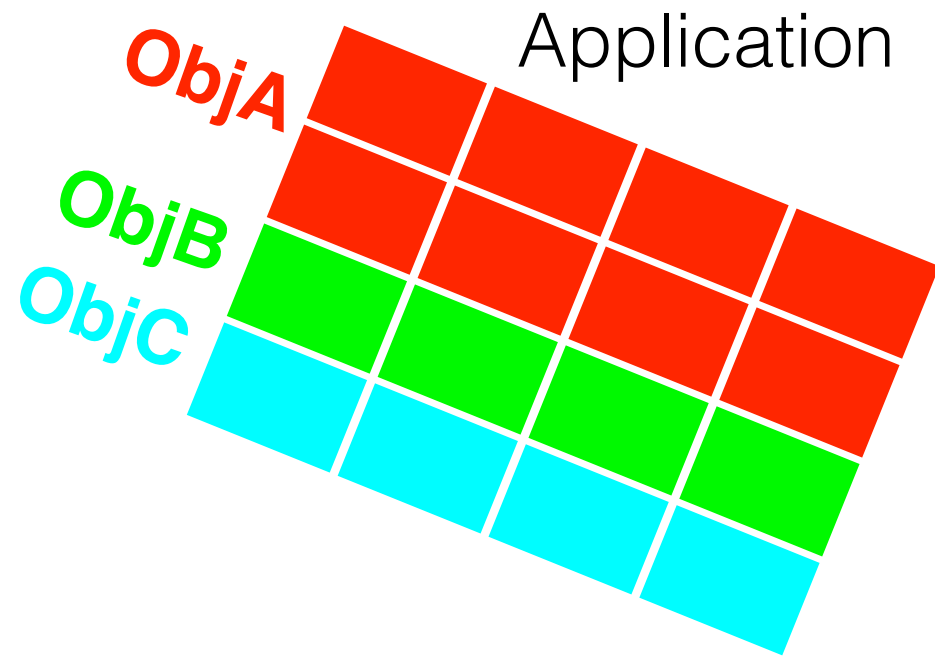
# Address→Object Mapping via Shadow Memory

Solution: shadow memory



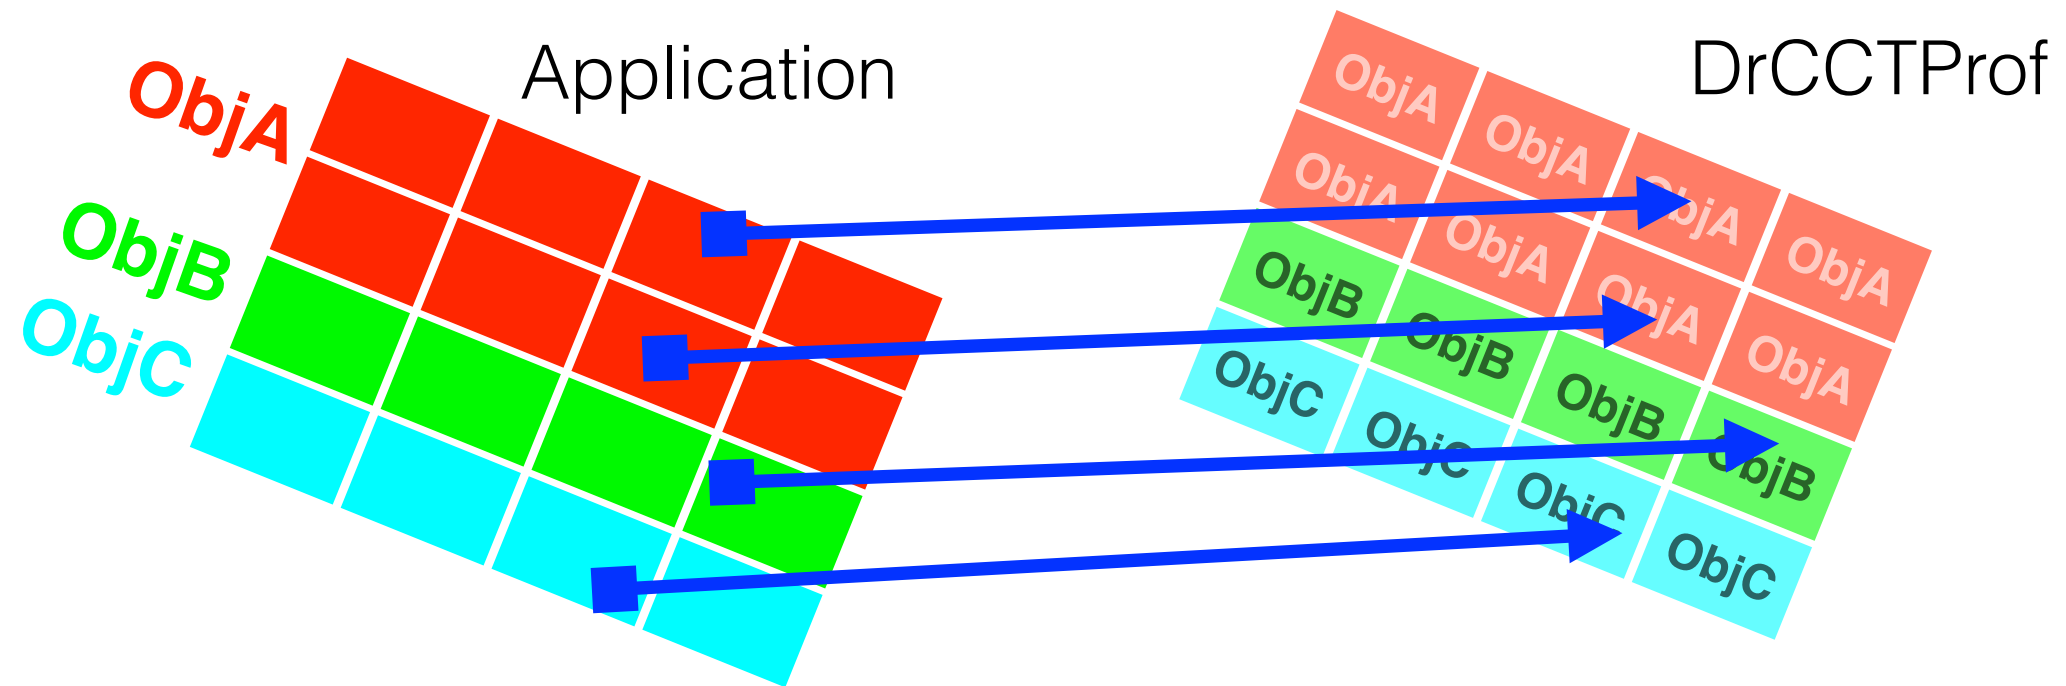
# Address→Object Mapping via Shadow Memory

Solution: shadow memory



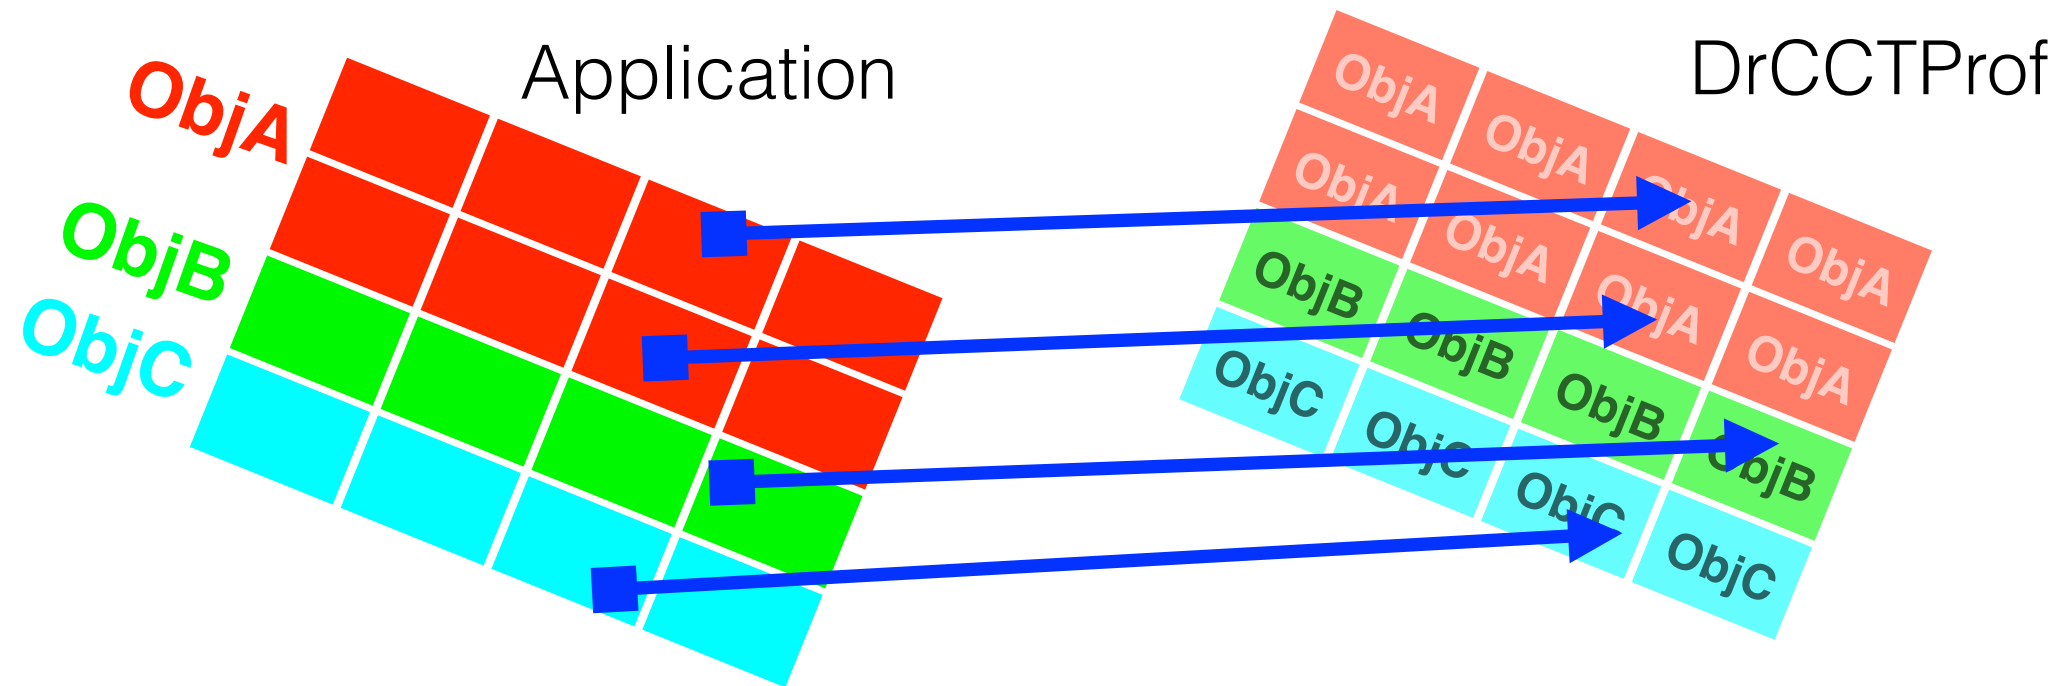
# Address→Object Mapping via Shadow Memory

Solution: shadow memory



# Address→Object Mapping via Shadow Memory

Solution: shadow memory



- For each memory cell, a shadow cell holds a handle for the memory cell's data object
$$O(\sum_{i=1}^N \text{sizeof}(\text{Obj}(i)))$$
  - ♦ Low lookup cost— $O(1)$ , high memory cost—
  - ♦ Shadow memory supports concurrent access

# Try DrCCTProf and Build Your Own Clients

- A number client tools available atop DrCCTProf
  - ♦ Check the example client tools at [DrCCTProf/src/clients](http://DrCCTProf/src/clients)
- You need to develop your own client tools to understand compiler code generation
  - ♦ Project 0: learn how to use DrCCTProf
    - \* Learn and write a simple but useful client
    - \* Understand compiler code generation

# Data Visualization

- Standard
  - ♦ Texts output: write human readable texts into a file
- Advanced
  - ♦ GUI-based visualization
    - \* Using HPCToolkit format
    - \* Using VSCode format

- Standard

- ♦ Texts c

- Advanced

- ♦ HPCTo

- \* Visual

The image shows a code editor window titled 'test\_app\_cct.cxx' with the following C++ code:

```

21 }
22
23 void t1_fun() {
24     for(int i = 0; i < 2222; i++){
25         t1_sub_fun();
26     }
27 }
28
29 void t2_fun() {
30     for(int i = 0; i < 1111; i++){
31         t2_sub_fun();
32     }
33 }
34 #endif
35
36 #ifdef MULTITHREADING
37 void *thread_1(void *arg)

```

Below the code editor is a call stack visualization window. It shows the current execution state with a table of scopes and instruction counts.

Scope	INS_COUNT
> <unknown procedure> 0x1093 [ld-2.27.so]	46276
▼ _start (test_app_cct)	43329
▼ > <unknown procedure> 0x21b95 [libc-2.27.so]	43329
▼ > main (test_app_cct)	43329
▼ > 82: t1_fun (test_app_cct)	28886
▼ loop at test_app_cct.cxx: 24	28886
▶ > 25: t1_sub_fun (test_app_cct)	17776
test_app_cct.cxx: 24	8888
<b>test_app_cct.cxx: 25</b>	<b>2222</b>
▼ > 83: t2_fun (test_app_cct)	14443
▼ loop at test_app_cct.cxx: 30	14443
▶ > 31: t2_sub_fun (test_app_cct)	8888
test_app_cct.cxx: 30	4444
test_app_cct.cxx: 31	1111

a file